

New York City Yellow Taxi Data

Objective

In this case study you will be learning exploratory data analysis (EDA) with the help of a dataset on yellow taxi rides in New York City. This will enable you to understand why EDA is an important step in the process of data science and machine learning.

Problem Statement

As an analyst at an upcoming taxi operation in NYC, you are tasked to use the 2023 taxi trip data to uncover insights that could help optimise taxi operations. The goal is to analyse patterns in the data that can inform strategic decisions to improve service efficiency, maximise revenue, and enhance passenger experience.

Tasks

You need to perform the following steps for successfully completing this assignment:

1. Data Loading
 2. Data Cleaning
 3. Exploratory Analysis: Bivariate and Multivariate
 4. Creating Visualisations to Support the Analysis
 5. Deriving Insights and Stating Conclusions
-

NOTE: The marks given along with headings and sub-headings are cumulative marks for those particular headings/sub-headings.

The actual marks for each task are specified within the tasks themselves.

For example, marks given with heading 2 or sub-heading 2.1 are the cumulative marks, for your reference only.

The marks you will receive for completing tasks are given with the tasks.

Suppose the marks for two tasks are: 3 marks for 2.1.1 and 2 marks for 3.2.2, or

- 2.1.1 [3 marks]
- 3.2.2 [2 marks]

then, you will earn 3 marks for completing task 2.1.1 and 2 marks for completing task 3.2.2.

Data Understanding

The yellow taxi trip records include fields capturing pick-up and drop-off dates/times, pick-up and drop-off locations, trip distances, itemized fares, rate types, payment types, and driver-reported passenger counts.

The data is stored in Parquet format (*.parquet*). The dataset is from 2009 to 2024. However, for this assignment, we will only be using the data from 2023.

The data for each month is present in a different parquet file. You will get twelve files for each of the months in 2023.

The data was collected and provided to the NYC Taxi and Limousine Commission (TLC) by technology providers like vendors and taxi hailing apps.

You can find the link to the TLC trip records page here: <https://www.nyc.gov/site/tlc/about/tlc-trip-record-data.page>

Data Description

You can find the data description here: [Data Dictionary](#)

Trip Records

Field Name	description
VendorID	A code indicating the TPEP provider that provided the record. 1= Creative Mobile Technologies, LLC; 2= VeriFone Inc.
tpep_pickup_datetime	The date and time when the meter was engaged.
tpep_dropoff_datetime	The date and time when the meter was disengaged.
Passenger_count	The number of passengers in the vehicle. This is a driver-entered value.
Trip_distance	The elapsed trip distance in miles reported by the taximeter.
PULocationID	TLC Taxi Zone in which the taximeter was engaged
DOLocationID	TLC Taxi Zone in which the taximeter was disengaged
RateCodeID	The final rate code in effect at the end of the trip. 1 = Standard rate 2 = JFK 3 = Newark 4 = Nassau or Westchester 5 = Negotiated fare 6 = Group ride
Store_and_fwd_flag	This flag indicates whether the trip

Field Name	description
	record was held in vehicle memory before sending to the vendor, aka "store and forward," because the vehicle did not have a connection to the server. Y= store and forward trip N= not a store and forward trip
Payment_type	A numeric code signifying how the passenger paid for the trip. 1 = Credit card 2 = Cash 3 = No charge 4 = Dispute 5 = Unknown 6 = Voided trip
Fare_amount	The time-and-distance fare calculated by the meter. Extra Miscellaneous extras and surcharges. Currently, this only includes the 0.50 and 1 USD rush hour and overnight charges.
MTA_tax	0.50 USD MTA tax that is automatically triggered based on the metered rate in use.
Improvement_surcharge	0.30 USD improvement surcharge assessed trips at the flag drop. The improvement surcharge began being levied in 2015.
Tip_amount	Tip amount – This field is automatically populated for credit card tips. Cash tips are not included.
Tolls_amount	Total amount of all tolls paid in trip.
total_amount	The total amount charged to passengers. Does not include cash tips.
Congestion_Surcharge	Total amount collected in trip for NYS congestion surcharge.
Airport_fee	1.25 USD for pick up only at LaGuardia and John F. Kennedy Airports

Although the amounts of extra charges and taxes applied are specified in the data dictionary, you will see that some cases have different values of these charges in the actual data.

Taxi Zones

Each of the trip records contains a field corresponding to the location of the pickup or drop-off of the trip, populated by numbers ranging from 1-263.

These numbers correspond to taxi zones, which may be downloaded as a table or map/shapefile and matched to the trip records using a join.

This is covered in more detail in later sections.

1 Data Preparation

[5 marks]

Import Libraries

```
# Import warnings

import warnings
warnings.filterwarnings("ignore")

# Import the libraries you will be using for analysis
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Recommended versions
# numpy version: 1.26.4
# pandas version: 2.2.2
# matplotlib version: 3.10.0
# seaborn version: 0.13.2

print("numpy version:", np.__version__)
print("pandas version:", pd.__version__)
print("matplotlib version:", plt.matplotlib.__version__)
print("seaborn version:", sns.__version__)

numpy version: 1.26.4
pandas version: 2.2.2
matplotlib version: 3.10.0
seaborn version: 0.13.2
```

1.1 Load the dataset

[5 marks]

You will see twelve files, one for each month.

To read parquet files with Pandas, you have to follow a similar syntax as that for CSV files.

```
df = pd.read_parquet('file.parquet')
```

```
# Try loading one file
```

```
df = pd.read_parquet('2023-1.parquet')
df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 3041714 entries, 0 to 3066765
Data columns (total 19 columns):
#   Column                                Dtype
---  -
0   VendorID                             int64
1   tpep_pickup_datetime                 datetime64[us]
2   tpep_dropoff_datetime               datetime64[us]
3   passenger_count                     float64
4   trip_distance                       float64
5   RatecodeID                         float64
6   store_and_fwd_flag                  object
7   PULocationID                       int64
8   DOLocationID                       int64
9   payment_type                       int64
10  fare_amount                         float64
11  extra                              float64
12  mta_tax                            float64
13  tip_amount                         float64
14  tolls_amount                       float64
15  improvement_surcharge              float64
16  total_amount                       float64
17  congestion_surcharge               float64
18  airport_fee                        float64
dtypes: datetime64[us](2), float64(12), int64(4), object(1)
memory usage: 464.1+ MB
```

How many rows are there? Do you think handling such a large number of rows is computationally feasible when we have to combine the data for all twelve months into one?

To handle this, we need to sample a fraction of data from each of the files. How to go about that? Think of a way to select only some portion of the data from each month's file that accurately represents the trends.

Sampling the Data

One way is to take a small percentage of entries for pickup in every hour of a date. So, for all the days in a month, we can iterate through the hours and select 5% values randomly from those. Use `tpep_pickup_datetime` for this. Separate date and hour from the datetime values and then for each date, select some fraction of trips for each of the 24 hours.

To sample data, you can use the `sample()` method. Follow this syntax:

```
# sampled_data is an empty DF to keep appending sampled data of each hour
# hour_data is the DF of entries for an hour 'X' on a date 'Y'
```

```

sample = hour_data.sample(frac = 0.05, random_state = 42)
# sample 0.05 of the hour_data
# random_state is just a seed for sampling, you can define it yourself

sampled_data = pd.concat([sampled_data, sample]) # adding data for
this hour to the DF

```

This *sampled_data* will contain 5% values selected at random from each hour.

Note that the code given above is only the part that will be used for sampling and not the complete code required for sampling and combining the data files.

Keep in mind that you sample by date AND hour, not just hour. (Why?)

1.1.1 [5 marks] Figure out how to sample and combine the files.

Note: It is not mandatory to use the method specified above. While sampling, you only need to make sure that your sampled data represents the overall data of all the months accurately.

```

# Sample the data
# It is recommended to not load all the files at once to avoid memory
overload

# from google.colab import drive
# drive.mount('/content/drive')

# Take a small percentage of entries from each hour of every date.
# Iterating through the monthly data:
#   read a month file -> day -> hour: append sampled data -> move to
next hour -> move to next day after 24 hours -> move to next month
file
# Create a single dataframe for the year combining all the monthly
data
import os

os.chdir(r'C:\Users\ADMIN\AL ML Coursesss\trip_records')

file_list = os.listdir()

df = pd.DataFrame()

for file_name in file_list:
    try:
        file_path = os.path.join(os.getcwd(), file_name)

```

```

        if file_name.endswith('.parquet'):
            sampled_data = pd.read_parquet(file_path,
engine='pyarrow')

            sampled_data = sampled_data.sample(frac=0.007,
random_state=42)

            df = pd.concat([df, sampled_data], ignore_index=True)

    except Exception as e:
        print(f"Error reading file {file_name}: {e}")

```

After combining the data files into one DataFrame, convert the new DataFrame to a CSV or parquet file and store it to use directly.

Ideally, you can try keeping the total entries to around 250,000 to 300,000.

```

# Store the df in csv/parquet
df.to_csv('data_New.csv', index=False)

```

2 Data Cleaning

[30 marks]

Now we can load the new data directly.

```

# Load the new data file
df = pd.read_csv('data_New.csv')

df.head()

```

	VendorID	passenger_count \	tpep_pickup_datetime	tpep_dropoff_datetime
0	1	2.0	2023-01-05 07:50:08	2023-01-05 08:02:04
1	2	5.0	2023-01-17 07:47:24	2023-01-17 08:00:50
2	2	1.0	2023-01-25 21:57:59	2023-01-25 22:00:33
3	2	2.0	2023-01-09 19:36:54	2023-01-09 19:52:01
4	1	1.0	2023-01-11 22:19:13	2023-01-11 22:32:37

	trip_distance	RatecodeID	store_and_fwd_flag	PULocationID
0	1.90	1.0	N	239

```

236
1      1.86      1.0      N      239
162
2      0.50      1.0      N      162
170
3      2.56      1.0      N      162
262
4      2.80      1.0      N      164
231

```

```

      payment_type  fare_amount  extra  mta_tax  tip_amount  tolls_amount
\
0      1      13.5      2.5      0.5      2.50      0.0
1      1      14.2      0.0      0.5      3.64      0.0
2      1      5.1      1.0      0.5      2.02      0.0
3      1      17.0      2.5      0.5      4.70      0.0
4      1      14.9      3.5      0.5      3.98      0.0

```

```

      improvement_surcharge  total_amount  congestion_surcharge
airport_fee \
0      1.0      20.00      2.5
0.0
1      1.0      21.84      2.5
0.0
2      1.0      12.12      2.5
0.0
3      1.0      28.20      2.5
0.0
4      1.0      23.88      2.5
0.0

```

```

      Airport_fee
0      NaN
1      NaN
2      NaN
3      NaN
4      NaN

```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 265500 entries, 0 to 265499
Data columns (total 20 columns):

```

```

#      Column      Non-Null Count  Dtype
---  -

```


0	VendorID	265500	non-null	int64
1	tpep_pickup_datetime	265500	non-null	object
2	tpep_dropoff_datetime	265500	non-null	object
3	passenger_count	256470	non-null	float64
4	trip_distance	265500	non-null	float64
5	RatecodeID	256470	non-null	float64
6	store_and_fwd_flag	256470	non-null	object
7	PULocationID	265500	non-null	int64
8	DOLocationID	265500	non-null	int64
9	payment_type	265500	non-null	int64
10	fare_amount	265500	non-null	float64
11	extra	265500	non-null	float64
12	mta_tax	265500	non-null	float64
13	tip_amount	265500	non-null	float64
14	tolls_amount	265500	non-null	float64
15	improvement_surcharge	265500	non-null	float64
16	total_amount	265500	non-null	float64
17	congestion_surcharge	256470	non-null	float64
18	airport_fee	20788	non-null	float64
19	Airport_fee	235682	non-null	float64

dtypes: float64(13), int64(4), object(3)
memory usage: 40.5+ MB

2.1 Fixing Columns

[10 marks]

Fix/drop any columns as you seem necessary in the below sections

2.1.1 [2 marks]

Fix the index and drop unnecessary columns

```
# Fix the index and drop any columns that are not needed
unnecessary_columns = [col for col in df.columns if "Unnamed" in col
or "index" in col.lower()]
df = df.drop(columns=unnecessary_columns, errors='ignore')
```

2.1.2 [3 marks] There are two airport fee columns. This is possibly an error in naming columns. Let's see whether these can be combined into a single column.

```
# Combine the two airport fee columns

if 'airport_fee' in df.columns and 'Airport_fee' in df.columns:
    df['Combined_Airport_Fee'] = df['airport_fee'].fillna(0) +
df['Airport_fee'].fillna(0)

df = df.drop(columns=['airport_fee', 'Airport_fee'])
```

```
df.head()
```

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime
0	1	2023-01-05 07:50:08	2023-01-05 08:02:04
1	2	2023-01-17 07:47:24	2023-01-17 08:00:50
2	2	2023-01-25 21:57:59	2023-01-25 22:00:33
3	2	2023-01-09 19:36:54	2023-01-09 19:52:01
4	1	2023-01-11 22:19:13	2023-01-11 22:32:37

	trip_distance	RatecodeID	store_and_fwd_flag	PULocationID
0	1.90	1.0	N	239
1	1.86	1.0	N	239
2	0.50	1.0	N	162
3	2.56	1.0	N	162
4	2.80	1.0	N	164

	payment_type	fare_amount	extra	mta_tax	tip_amount	tolls_amount
0	1	13.5	2.5	0.5	2.50	0.0
1	1	14.2	0.0	0.5	3.64	0.0
2	1	5.1	1.0	0.5	2.02	0.0
3	1	17.0	2.5	0.5	4.70	0.0
4	1	14.9	3.5	0.5	3.98	0.0

	improvement_surcharge	total_amount	congestion_surcharge
0	1.0	20.00	2.5
1	1.0	21.84	2.5
2	1.0	12.12	2.5
3	1.0	28.20	2.5
4	1.0	23.88	2.5

Combined_Airport_Fee

0	0.0
1	0.0
2	0.0
3	0.0
4	0.0

2.1.4 [5 marks] Fix columns with negative (monetary) values

```
# check where values of fare amount are negative
print(df["fare_amount"].min())

# No negative value
0.0
```

Did you notice something different in the `RatecodeID` column for above records?

Yes we have found 99 in the record

```
# Analyse RatecodeID for the negative fare amounts
df['RatecodeID'].value_counts()

RatecodeID
1.0    241967
2.0    10228
5.0     1488
99.0    1452
3.0     829
4.0     506
Name: count, dtype: int64

# Find which columns have negative values
df = df[df["RatecodeID"] != 99.0]
print(df["RatecodeID"].value_counts())

RatecodeID
1.0    241967
2.0    10228
5.0     1488
3.0     829
4.0     506
Name: count, dtype: int64

df = df[df["RatecodeID"] != 99.0]
print(df["RatecodeID"].value_counts())

RatecodeID
1.0    241967
2.0    10228
```

```
5.0      1488
3.0       829
4.0       506
```

```
Name: count, dtype: int64
```

```
# fix these negative values
```

```
print(df.describe())
```

	VendorID	passenger_count	trip_distance	RatecodeID	\
count	264048.000000	255018.000000	264048.000000	255018.000000	
mean	1.741608	1.374350	3.638318	1.075901	
std	0.442827	0.899321	80.015295	0.398203	
min	1.000000	0.000000	0.000000	1.000000	
25%	1.000000	1.000000	1.040000	1.000000	
50%	2.000000	1.000000	1.780000	1.000000	
75%	2.000000	1.000000	3.380000	1.000000	
max	6.000000	8.000000	37523.740000	5.000000	

	PULocationID	DOLocationID	payment_type	fare_amount	\
count	264048.000000	264048.000000	264048.000000	264048.000000	
mean	165.313057	164.200104	1.165224	19.779402	
std	63.884569	69.654169	0.510045	18.308347	
min	1.000000	1.000000	0.000000	0.000000	
25%	132.000000	114.000000	1.000000	9.300000	
50%	162.000000	162.000000	1.000000	13.500000	
75%	234.000000	234.000000	1.000000	21.900000	
max	265.000000	265.000000	4.000000	750.000000	

	extra	mta_tax	tip_amount	tolls_amount	\
count	264048.000000	264048.000000	264048.000000	264048.000000	
mean	1.586107	0.495319	3.577066	0.591301	
std	1.828371	0.048421	4.085435	2.170130	
min	-1.000000	-0.500000	0.000000	0.000000	
25%	0.000000	0.500000	1.000000	0.000000	
50%	1.000000	0.500000	2.860000	0.000000	
75%	2.500000	0.500000	4.450000	0.000000	
max	13.000000	0.800000	288.000000	132.040000	

	improvement_surcharge	total_amount	congestion_surcharge	\
count	264048.000000	264048.000000	255018.000000	
mean	0.999140	28.890812	2.323542	
std	0.029012	22.919192	0.640664	
min	-1.000000	-5.000000	-2.500000	
25%	1.000000	15.960000	2.500000	
50%	1.000000	21.000000	2.500000	
75%	1.000000	30.720000	2.500000	
max	1.000000	757.940000	2.500000	

	Combined_Airport_Fee
count	264048.000000

```

mean      0.140314
std       0.461338
min       -1.250000
25%       0.000000
50%       0.000000
75%       0.000000
max       1.750000

```

fix these negative values

```

monetary_columns = ["extra", "mta_tax", "improvement_surcharge",
                    "total_amount", "congestion_surcharge",
                    "Combined_Airport_Fee"]

```

```

df[monetary_columns] = df[monetary_columns].abs()

```

```

print(df.describe())

```

	VendorID	passenger_count	trip_distance	RatecodeID	\
count	264048.000000	255018.000000	264048.000000	255018.000000	
mean	1.741608	1.374350	3.638318	1.075901	
std	0.442827	0.899321	80.015295	0.398203	
min	1.000000	0.000000	0.000000	1.000000	
25%	1.000000	1.000000	1.040000	1.000000	
50%	2.000000	1.000000	1.780000	1.000000	
75%	2.000000	1.000000	3.380000	1.000000	
max	6.000000	8.000000	37523.740000	5.000000	

	PULocationID	DOLocationID	payment_type	fare_amount	\
count	264048.000000	264048.000000	264048.000000	264048.000000	
mean	165.313057	164.200104	1.165224	19.779402	
std	63.884569	69.654169	0.510045	18.308347	
min	1.000000	1.000000	0.000000	0.000000	
25%	132.000000	114.000000	1.000000	9.300000	
50%	162.000000	162.000000	1.000000	13.500000	
75%	234.000000	234.000000	1.000000	21.900000	
max	265.000000	265.000000	4.000000	750.000000	

	extra	mta_tax	tip_amount	tolls_amount	\
count	264048.000000	264048.000000	264048.000000	264048.000000	
mean	1.586114	0.495364	3.577066	0.591301	
std	1.828364	0.047954	4.085435	2.170130	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.500000	1.000000	0.000000	
50%	1.000000	0.500000	2.860000	0.000000	
75%	2.500000	0.500000	4.450000	0.000000	
max	13.000000	0.800000	288.000000	132.040000	

	improvement_surcharge	total_amount	congestion_surcharge	\
count	264048.000000	264048.000000	255018.000000	
mean	0.999246	28.891161	2.323718	

std	0.025095	22.918753	0.640024
min	0.000000	0.000000	0.000000
25%	1.000000	15.960000	2.500000
50%	1.000000	21.000000	2.500000
75%	1.000000	30.720000	2.500000
max	1.000000	757.940000	2.500000

	Combined_Airport_Fee
count	264048.000000
mean	0.140333
std	0.461332
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.750000

2.2 Handling Missing Values

[10 marks]

2.2.1 [2 marks] Find the proportion of missing values in each column

Find the proportion of missing values in each column

```
df.isnull().sum()
```

```
VendorID      0
tpep_pickup_datetime      0
tpep_dropoff_datetime     0
passenger_count    9030
trip_distance      0
RatecodeID      9030
store_and_fwd_flag    9030
PULocationID     0
DOLocationID     0
payment_type      0
fare_amount      0
extra            0
mta_tax          0
tip_amount       0
tolls_amount     0
improvement_surcharge    0
total_amount      0
congestion_surcharge    9030
Combined_Airport_Fee      0
dtype: int64
```

2.2.2 [3 marks] Handling missing values in `passenger_count`

```
# Display the rows with null values
```

```
df[df['passenger_count'].isnull()]
```

```
# Impute NaN values in 'passenger_count'
```

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime
passenger_count \			
155	2	2023-01-20 15:15:17	2023-01-20 15:27:06
NaN			
157	2	2023-01-21 15:17:59	2023-01-21 15:39:36
NaN			
303	2	2023-01-28 23:58:47	2023-01-29 00:11:44
NaN			
309	1	2023-01-21 02:51:57	2023-01-21 03:15:22
NaN			
356	2	2023-01-17 21:22:28	2023-01-17 21:36:25
NaN			
...
...			
265403	1	2023-09-13 22:16:27	2023-09-13 22:29:21
NaN			
265421	1	2023-09-11 21:23:50	2023-09-11 21:35:41
NaN			
265445	2	2023-09-19 21:24:31	2023-09-19 21:41:47
NaN			
265448	1	2023-09-08 22:38:30	2023-09-08 22:50:37
NaN			
265477	2	2023-09-12 22:04:59	2023-09-12 22:21:59
NaN			

	trip_distance	RatecodeID	store_and_fwd_flag	PULocationID
155	1.35	NaN	NaN	142
157	6.67	NaN	NaN	170
303	1.73	NaN	NaN	211
309	0.00	NaN	NaN	79
356	2.43	NaN	NaN	143
...
265403	0.00	NaN	NaN	233
265421	0.00	NaN	NaN	141
265445	4.00	NaN	NaN	148
265448	0.00	NaN	NaN	249
265477	4.23	NaN	NaN	161

	DOLocationID	payment_type	fare_amount	extra	mta_tax
tip_amount \					
155	239	0	13.65	0.0	0.5
3.53					
157	41	0	27.71	0.0	0.5
7.34					

303	4	0	15.55	0.0	0.5
3.91					
309	166	0	22.83	0.0	0.5
0.00					
356	263	0	15.21	0.0	0.5
2.88					
...
...					
265403	90	0	17.11	0.0	0.5
0.00					
265421	107	0	17.50	0.0	0.5
0.00					
265445	181	0	21.23	0.0	0.5
5.05					
265448	164	0	42.92	0.0	0.5
0.00					
265477	74	0	23.54	0.0	0.5
0.00					

	tolls_amount	improvement_surcharge	total_amount \
155	0.0	1.0	21.18
157	0.0	1.0	39.05
303	0.0	1.0	23.46
309	0.0	1.0	26.83
356	0.0	1.0	22.09
...
265403	0.0	1.0	21.11
265421	0.0	1.0	21.50
265445	0.0	1.0	30.28
265448	0.0	1.0	46.92
265477	0.0	1.0	27.54

	congestion_surcharge	Combined_Airport_Fee
155	NaN	0.0
157	NaN	0.0
303	NaN	0.0
309	NaN	0.0
356	NaN	0.0
...
265403	NaN	0.0
265421	NaN	0.0
265445	NaN	0.0
265448	NaN	0.0
265477	NaN	0.0

[9030 rows x 19 columns]

```
# Impute NaN values in 'passenger_count'
```

```
df["passenger_count"] =
```



```
df["passenger_count"].fillna(df["passenger_count"].median())
df['passenger_count'].isnull().value_counts()

passenger_count
False      264048
Name: count, dtype: int64
```

Did you find zeroes in passenger_count? Handle these.

Yes first i could zero value in Passenger_count now i have removed it.

```
df['passenger_count'].value_counts()

df = df[df["passenger_count"] > 0]

df['passenger_count'].value_counts()

passenger_count
1.0      200322
2.0      38901
3.0       9625
4.0       5426
5.0       3459
6.0       2250
8.0         2
Name: count, dtype: int64
```

2.2.3 [2 marks] Handle missing values in RatecodeID

```
# Fix missing values in 'RatecodeID'
df['RatecodeID'].isnull().value_counts()

df["RatecodeID"] = df["RatecodeID"].fillna(df["RatecodeID"].median())

df['RatecodeID'].isnull().value_counts()

RatecodeID
False      259985
Name: count, dtype: int64
```

2.2.4 [3 marks] Impute NaN in congestion_surcharge

```
# handle null values in congestion_surcharge

df["congestion_surcharge"] =
df["congestion_surcharge"].fillna(df["congestion_surcharge"].median())
df['congestion_surcharge'].isnull().value_counts()
```

```
congestion_surcharge
False      259985
Name: count, dtype: int64
```

Are there missing values in other columns? Did you find NaN values in some other set of columns? Handle those missing values below.

```
# Handle any remaining missing values
```

```
df.isnull().sum()
```

```
VendorID      0
tpep_pickup_datetime  0
tpep_dropoff_datetime  0
passenger_count  0
trip_distance  0
RatecodeID     0
store_and_fwd_flag    9030
PULocationID    0
DOLocationID    0
payment_type     0
fare_amount     0
extra           0
mta_tax         0
tip_amount      0
tolls_amount    0
improvement_surcharge  0
total_amount    0
congestion_surcharge  0
Combined_Airport_Fee  0
dtype: int64
```

```
df["store_and_fwd_flag"] = df["store_and_fwd_flag"].fillna('N')
df['store_and_fwd_flag'].isnull().value_counts()
```

```
store_and_fwd_flag
False      259985
Name: count, dtype: int64
```

2.3 Handling Outliers

[10 marks]

Before we start fixing outliers, let's perform outlier analysis.

```
# Describe the data and check if there are any potential outliers present
# Check for potential out of place values in various columns
```

2.3.1 [10 marks] Based on the above analysis, it seems that some of the outliers are present due to errors in registering the trips. Fix the outliers.

Some points you can look for:


- Entries where `trip_distance` is nearly 0 and `fare_amount` is more than 300
- Entries where `trip_distance` and `fare_amount` are 0 but the pickup and dropoff zones are different (both distance and fare should not be zero for different zones)
- Entries where `trip_distance` is more than 250 miles.
- Entries where `payment_type` is 0 (there is no `payment_type` 0 defined in the data dictionary)

These are just some suggestions. You can handle outliers in any way you wish, using the insights from above outlier analysis.

How will you fix each of these values? Which ones will you drop and which ones will you replace?

First, let us remove 7+ passenger counts as there are very less instances.

```
# remove passenger_count > 6
df = df[df["passenger_count"] <= 7]
plt.boxplot(df['passenger_count'])
plt.show()
```



```
# Continue with outlier handling
#Entries where `trip_distance` is nearly 0 and `fare_amount` is more
than 300
```

```
filtered_df = df[(df['trip_distance'] < 0.1) & (df['fare_amount'] >
300)]
print(filtered_df)
```

```
df = df[~((df["trip_distance"] < 0.1) & (df["fare_amount"] > 300))]
```

```
filtered_df = df[(df['trip_distance'] < 0.1) & (df['fare_amount'] >
300)]
print(filtered_df)
```

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime
passenger_count \			
47758	2	2023-11-26 16:04:06	2023-11-26 16:04:12
1.0			
109684	2	2023-03-19 23:12:29	2023-03-19 23:12:35
1.0			
210866	2	2023-05-16 19:12:48	2023-05-16 19:12:51
1.0			
226046	2	2023-07-24 21:43:56	2023-07-24 21:57:57
1.0			
235276	1	2023-07-20 16:17:03	2023-07-20 16:17:23
1.0			

	trip_distance	RatecodeID	store_and_fwd_flag	PULocationID	\
47758	0.0	5.0	N	265	
109684	0.0	5.0	N	265	
210866	0.0	5.0	N	265	
226046	0.0	5.0	N	265	
235276	0.0	5.0	N	193	

	DOLocationID	payment_type	fare_amount	extra	mta_tax
tip_amount \					
47758	265	1	305.14	0.0	0.0
61.23					
109684	265	2	533.00	0.0	0.0
0.00					
210866	265	1	396.00	0.0	0.5
0.80					
226046	265	2	500.00	0.0	0.0
0.00					
235276	193	2	555.54	0.0	0.0
0.00					

	tolls_amount	improvement_surcharge	total_amount	\
47758	0.0	1.0	367.37	
109684	0.0	1.0	534.00	

210866	0.0	1.0	398.30
226046	0.0	1.0	501.00
235276	0.0	1.0	556.54

	congestion_surcharge	Combined_Airport_Fee
47758	0.0	0.0
109684	0.0	0.0
210866	0.0	0.0
226046	0.0	0.0
235276	0.0	0.0

Empty DataFrame

Columns: [VendorID, tpep_pickup_datetime, tpep_dropoff_datetime, passenger_count, trip_distance, RatecodeID, store_and_fwd_flag, PULocationID, DOLocationID, payment_type, fare_amount, extra, mta_tax, tip_amount, tolls_amount, improvement_surcharge, total_amount, congestion_surcharge, Combined_Airport_Fee]

Index: []

#Entries where `trip_distance` and `fare_amount` are 0 but the pickup and dropoff zones are different (both distance and fare should not be zero for different zones)

```
filtered_df1 = df[(df['trip_distance'] < 0.1) & (df['fare_amount'] < 0.01)]
print(filtered_df1)
```

```
df= df[~((df['trip_distance'] < 0.1) & (df['fare_amount'] < 0.01))]
```

```
filtered_df1 = df[(df['trip_distance'] < 0.1) & (df['fare_amount'] < 0.01)]
print(filtered_df1)
```

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime
passenger_count \			
3835	1	2023-01-24 18:31:57	2023-01-24 18:33:16
2.0			
6585	2	2023-01-13 06:19:23	2023-01-13 06:19:27
1.0			
9570	2	2023-01-14 12:03:42	2023-01-14 12:03:49
2.0			
9649	1	2023-01-29 09:52:41	2023-01-29 09:53:54
1.0			
11054	2	2023-01-19 13:28:25	2023-01-19 13:28:51
1.0			
15245	2	2023-01-12 12:34:27	2023-01-12 12:35:28
1.0			
19126	1	2023-01-04 13:09:40	2023-01-04 13:09:56
1.0			
23206	1	2023-10-26 17:20:38	2023-10-26 17:21:18
3.0			

23897 2.0	2	2023-10-07 11:51:21	2023-10-07 11:51:30
31754 1.0	2	2023-10-19 12:31:08	2023-10-19 12:31:24
37179 1.0	2	2023-10-01 15:43:56	2023-10-01 15:45:51
57266 2.0	2	2023-11-18 23:45:16	2023-11-18 23:45:38
61078 1.0	2	2023-11-04 18:07:19	2023-11-04 18:08:02
76548 1.0	2	2023-12-12 07:22:35	2023-12-12 07:22:51
84134 1.0	1	2023-12-16 16:35:24	2023-12-16 16:35:24
85711 1.0	2	2023-12-15 10:48:42	2023-12-15 10:52:12
101617 1.0	1	2023-03-01 15:51:07	2023-03-01 15:51:27
103635 1.0	1	2023-03-21 08:10:43	2023-03-21 08:11:05
110247 1.0	2	2023-03-08 10:48:46	2023-03-08 10:51:24
120079 1.0	1	2023-06-01 14:21:00	2023-06-01 14:21:49
122049 4.0	2	2023-06-04 09:46:55	2023-06-04 09:47:05
140003 1.0	2	2023-08-11 16:19:41	2023-08-11 16:20:14
140992 2.0	1	2023-08-21 05:47:49	2023-08-21 05:47:49
141216 1.0	2	2023-08-31 20:27:23	2023-08-31 20:28:20
160098 2.0	1	2023-02-18 17:49:17	2023-02-18 17:49:30
161037 1.0	1	2023-02-10 00:42:04	2023-02-10 00:42:04
162406 1.0	1	2023-02-22 10:33:21	2023-02-22 10:33:40
169469 1.0	2	2023-02-25 10:49:04	2023-02-25 10:49:12
183715 1.0	1	2023-04-17 06:32:54	2023-04-17 06:33:35
189099 1.0	2	2023-04-24 19:16:53	2023-04-24 19:16:58
189345 2.0	1	2023-04-14 00:36:54	2023-04-14 00:36:54
189954 1.0	1	2023-04-17 12:45:58	2023-04-17 12:46:29
194062	2	2023-04-29 10:24:56	2023-04-29 10:26:26

1.0					
206373	2	2023-05-04	11:56:27	2023-05-04	11:56:37
1.0					
209872	2	2023-05-04	18:40:47	2023-05-04	18:40:55
1.0					
212354	2	2023-05-29	16:38:17	2023-05-29	16:38:27
2.0					
224069	1	2023-05-29	06:59:22	2023-05-29	06:59:22
1.0					
224660	2	2023-05-23	13:49:46	2023-05-23	13:50:01
1.0					
234318	1	2023-07-13	17:40:21	2023-07-13	17:41:47
1.0					
246157	2	2023-09-27	17:02:30	2023-09-27	17:08:36
1.0					
250671	1	2023-09-06	00:23:52	2023-09-06	00:25:30
1.0					

	trip_distance	RatecodeID	store_and_fwd_flag	PULocationID	\
3835	0.00	2.0	Y	237	
6585	0.00	1.0	N	132	
9570	0.00	2.0	N	13	
9649	0.00	5.0	N	188	
11054	0.00	1.0	N	193	
15245	0.00	1.0	N	264	
19126	0.00	5.0	N	132	
23206	0.00	4.0	N	7	
23897	0.01	2.0	N	238	
31754	0.00	1.0	N	264	
37179	0.00	1.0	N	193	
57266	0.00	1.0	N	162	
61078	0.00	1.0	N	246	
76548	0.00	1.0	N	239	
84134	0.00	5.0	Y	163	
85711	0.04	2.0	N	261	
101617	0.00	5.0	N	132	
103635	0.00	1.0	N	237	
110247	0.00	1.0	N	193	
120079	0.00	1.0	N	100	
122049	0.00	5.0	N	265	
140003	0.00	1.0	N	264	
140992	0.00	3.0	Y	162	
141216	0.01	1.0	N	264	
160098	0.00	1.0	N	50	
161037	0.00	1.0	N	26	
162406	0.00	5.0	N	68	
169469	0.00	1.0	N	264	
183715	0.00	1.0	N	89	
189099	0.00	1.0	N	132	

189345	0.00	5.0	Y	90
189954	0.00	5.0	N	145
194062	0.03	1.0	N	230
206373	0.00	1.0	N	193
209872	0.00	2.0	N	170
212354	0.00	1.0	N	68
224069	0.00	2.0	Y	132
224660	0.00	2.0	N	239
234318	0.00	5.0	N	132
246157	0.00	5.0	N	246
250671	0.00	5.0	N	107

	D0LocationID	payment_type	fare_amount	extra	mta_tax
tip_amount \					
3835	237	3	0.0	0.00	0.0
0.0					
6585	132	2	0.0	0.00	0.5
0.0					
9570	13	2	0.0	0.00	0.5
0.0					
9649	188	3	0.0	0.00	0.0
0.0					
11054	193	1	0.0	0.00	0.0
0.0					
15245	264	1	0.0	0.00	0.0
0.0					
19126	132	2	0.0	1.25	0.0
0.0					
23206	7	3	0.0	2.50	0.5
0.0					
23897	238	2	0.0	0.00	0.5
0.0					
31754	264	2	0.0	0.00	0.0
0.0					
37179	193	1	0.0	0.00	0.0
0.0					
57266	162	3	0.0	1.00	0.5
0.0					
61078	246	2	0.0	0.00	0.5
0.0					
76548	239	2	0.0	0.00	0.5
0.0					
84134	264	2	0.0	0.00	0.0
0.0					
85711	261	2	0.0	0.00	0.5
0.0					
101617	132	3	0.0	1.25	0.0
0.0					
103635	237	3	0.0	0.00	0.0

0.0					
110247	193	1	0.0	0.00	0.0
0.0					
120079	100	3	0.0	0.00	0.0
0.0					
122049	265	2	0.0	0.00	0.0
0.0					
140003	264	1	0.0	0.00	0.0
0.0					
140992	264	3	0.0	0.00	0.0
0.0					
141216	264	2	0.0	0.00	0.5
0.0					
160098	50	4	0.0	0.00	0.0
0.0					
161037	26	1	0.0	0.00	0.0
0.0					
162406	68	2	0.0	0.00	0.0
0.0					
169469	264	1	0.0	0.00	0.0
0.9					
183715	89	1	0.0	0.00	0.0
0.0					
189099	132	2	0.0	0.00	0.5
0.0					
189345	264	2	0.0	0.00	0.0
0.0					
189954	145	1	0.0	0.00	0.0
0.0					
194062	230	2	0.0	0.00	0.5
0.0					
206373	193	1	0.0	0.00	0.0
0.0					
209872	170	2	0.0	0.00	0.5
0.0					
212354	68	2	0.0	0.00	0.5
0.0					
224069	264	3	0.0	0.00	0.0
0.0					
224660	239	2	0.0	0.00	0.5
0.0					
234318	132	2	0.0	0.00	0.0
0.0					
246157	246	2	0.0	0.00	0.0
0.0					
250671	107	3	0.0	0.00	0.0
0.0					

tolls_amount	improvement_surcharge	total_amount	\
--------------	-----------------------	--------------	---

3835	0.0	0.0	0.00
6585	0.0	1.0	2.75
9570	0.0	1.0	4.00
9649	0.0	1.0	1.00
11054	0.0	0.0	0.00
15245	0.0	0.0	0.00
19126	0.0	1.0	2.25
23206	0.0	1.0	4.00
23897	0.0	1.0	4.00
31754	0.0	0.0	0.00
37179	0.0	0.0	0.00
57266	0.0	1.0	5.00
61078	0.0	1.0	4.00
76548	0.0	1.0	4.00
84134	0.0	0.0	0.00
85711	0.0	1.0	4.00
101617	0.0	1.0	2.25
103635	0.0	0.0	0.00
110247	0.0	0.0	0.00
120079	0.0	0.0	0.00
122049	0.0	1.0	1.00
140003	0.0	0.0	0.00
140992	0.0	0.0	0.00
141216	0.0	1.0	4.00
160098	0.0	0.0	0.00
161037	0.0	0.0	0.00
162406	0.0	1.0	1.00
169469	0.0	0.0	0.90
183715	0.0	0.0	0.00
189099	0.0	1.0	1.50
189345	0.0	0.0	0.00
189954	0.0	0.0	0.00
194062	0.0	1.0	4.00
206373	0.0	0.0	0.00
209872	0.0	1.0	4.00
212354	0.0	1.0	4.00
224069	0.0	0.0	0.00
224660	0.0	1.0	4.00
234318	0.0	1.0	1.00
246157	0.0	1.0	3.50
250671	0.0	1.0	1.00

	congestion_surcharge	Combined_Airport_Fee
3835	0.0	0.00
6585	0.0	1.25
9570	2.5	0.00
9649	0.0	0.00
11054	0.0	0.00
15245	0.0	0.00

19126	0.0	1.25
23206	0.0	0.00
23897	2.5	0.00
31754	0.0	0.00
37179	0.0	0.00
57266	2.5	0.00
61078	2.5	0.00
76548	2.5	0.00
84134	0.0	0.00
85711	2.5	0.00
101617	0.0	1.25
103635	0.0	0.00
110247	0.0	0.00
120079	0.0	0.00
122049	0.0	0.00
140003	0.0	0.00
140992	0.0	0.00
141216	2.5	0.00
160098	0.0	0.00
161037	0.0	0.00
162406	0.0	0.00
169469	0.0	0.00
183715	0.0	0.00
189099	0.0	0.00
189345	0.0	0.00
189954	0.0	0.00
194062	2.5	0.00
206373	0.0	0.00
209872	2.5	0.00
212354	2.5	0.00
224069	0.0	0.00
224660	2.5	0.00
234318	0.0	0.00
246157	2.5	0.00
250671	0.0	0.00


Empty DataFrame

Columns: [VendorID, tpep_pickup_datetime, tpep_dropoff_datetime, passenger_count, trip_distance, RatecodeID, store_and_fwd_flag, PULocationID, DOLocationID, payment_type, fare_amount, extra, mta_tax, tip_amount, tolls_amount, improvement_surcharge, total_amount, congestion_surcharge, Combined_Airport_Fee]


Index: []

#Entries where `trip_distance` is more than 250 miles.

```
plt.boxplot(df['trip_distance'])
plt.show()
```



```
df = df[~(df["trip_distance"] > 250)]  
plt.boxplot(df['trip_distance'])  
plt.show()
```



#Entries where `payment_type` is 0 (there is no payment_type 0 defined in the data dictionary)

```
df['payment_type'].value_counts()
```

```
payment_type
```

```
1    204312
```

```
2    43539
```

```
0     9028
```

```
4     1930
```

```
3     1126
```

```
Name: count, dtype: int64
```

#Entries where `payment_type` is 0 (there is no payment_type 0 defined in the data dictionary)

```
df = df[df["payment_type"] != 0]
```

```
df['payment_type'].value_counts()
```

```
payment_type
```

```
1    204312
```

```
2    43539
```

```
4     1930
```

```
3     1126
```

```
Name: count, dtype: int64
```

Do any columns need standardising?

```
print(df.describe())
```

	VendorID	passenger_count	trip_distance	RatecodeID	\
count	250907.000000	250907.000000	250907.000000	250907.000000	
mean	1.757583	1.396569	3.444610	1.075733	
std	0.428546	0.889117	4.586306	0.396638	
min	1.000000	1.000000	0.000000	1.000000	
25%	2.000000	1.000000	1.050000	1.000000	
50%	2.000000	1.000000	1.780000	1.000000	
75%	2.000000	1.000000	3.350000	1.000000	
max	2.000000	6.000000	170.300000	5.000000	

	PULocationID	DOLocationID	payment_type	fare_amount	\
count	250907.000000	250907.000000	250907.000000	250907.000000	
mean	165.403125	164.414807	1.205578	19.721341	
std	63.575002	69.610189	0.467381	18.367818	
min	1.000000	1.000000	1.000000	0.000000	
25%	132.000000	114.000000	1.000000	9.300000	
50%	162.000000	162.000000	1.000000	13.500000	
75%	234.000000	234.000000	1.000000	21.900000	
max	265.000000	265.000000	4.000000	750.000000	

	extra	mta_tax	tip_amount	tolls_amount	\
count	250907.000000	250907.000000	250907.000000	250907.000000	
mean	1.609131	0.495418	3.613894	0.594723	
std	1.826910	0.047678	4.112515	2.176181	
min	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	0.500000	1.000000	0.000000	
50%	1.000000	0.500000	2.860000	0.000000	
75%	2.500000	0.500000	4.480000	0.000000	
max	11.750000	0.800000	288.000000	132.040000	

	improvement_surcharge	total_amount	congestion_surcharge	\
count	250907.000000	250907.000000	250907.000000	
mean	0.999546	28.911278	2.323819	
std	0.019357	23.054619	0.639855	
min	0.000000	0.000000	0.000000	
25%	1.000000	15.960000	2.500000	
50%	1.000000	21.000000	2.500000	
75%	1.000000	30.620000	2.500000	
max	1.000000	757.940000	2.500000	

	Combined_Airport_Fee
count	250907.000000
mean	0.146118
std	0.469900
min	0.000000
25%	0.000000

50%	0.000000
75%	0.000000
max	1.750000

3 Exploratory Data Analysis

[90 marks]

```
df.columns.tolist()

['VendorID',
 'tpep_pickup_datetime',
 'tpep_dropoff_datetime',
 'passenger_count',
 'trip_distance',
 'RatecodeID',
 'store_and_fwd_flag',
 'PULocationID',
 'DOLocationID',
 'payment_type',
 'fare_amount',
 'extra',
 'mta_tax',
 'tip_amount',
 'tolls_amount',
 'improvement_surcharge',
 'total_amount',
 'congestion_surcharge',
 'Combined_Airport_Fee']
```

3.1 General EDA: Finding Patterns and Trends

[40 marks]

3.1.1 [3 marks] Categorise the variables into Numerical or Categorical..

- VendorID:
- tpep_pickup_datetime:
- tpep_dropoff_datetime:
- passenger_count:
- trip_distance:
- RatecodeID:
- PULocationID:
- DOLocationID:
- payment_type:
- pickup_hour:
- trip_duration:

The following monetary parameters belong in the same category, is it categorical or numerical?

- fare_amount
- extra
- mta_tax
- tip_amount
- tolls_amount
- improvement_surcharge
- total_amount
- congestion_surcharge
- airport_fee

Temporal Analysis

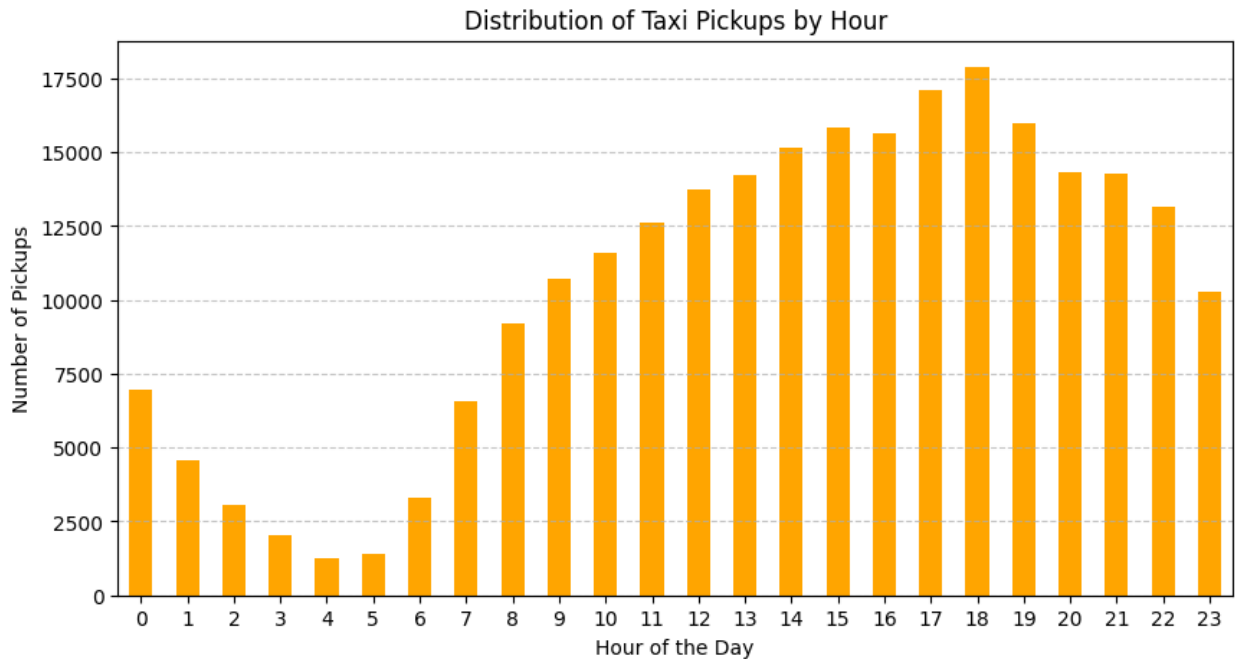
3.1.2 [5 marks] Analyse the distribution of taxi pickups by hours, days of the week, and months.

```
# Find and show the hourly trends in taxi pickups

df["tpep_pickup_datetime"] =
pd.to_datetime(df["tpep_pickup_datetime"])


df["pickup_hour"] = df["tpep_pickup_datetime"].dt.hour

plt.figure(figsize=(10, 5))
df["pickup_hour"].value_counts().sort_index().plot(kind="bar",
color="orange")
plt.xlabel("Hour of the Day")
plt.ylabel("Number of Pickups")
plt.title("Distribution of Taxi Pickups by Hour")
plt.xticks(rotation=0)
plt.grid(axis="y", linestyle="--", alpha=0.7)
plt.show()
```


```
# Find and show the daily trends in taxi pickups (days of the week)

df["pickup_day"] = df["tpep_pickup_datetime"].dt.dayofweek
plt.figure(figsize=(10, 5))
df["pickup_day"].value_counts().sort_index().plot(kind="bar", color="orange")
plt.xlabel("Day of the Week (0=Monday, 6=Sunday)")
plt.ylabel("Number of Pickups")
plt.title("Distribution of Taxi Pickups by Day of the Week")
plt.xticks(ticks=range(7), labels=["Mon", "Tue", "Wed", "Thu", "Fri", "Sat", "Sun"], rotation=0)
plt.grid(axis="y", linestyle="--", alpha=0.7)
plt.show()
```



```
# Show the monthly trends in pickups
```

```
df["tpep_pickup_datetime"] =  
pd.to_datetime(df["tpep_pickup_datetime"])  
  
df["pickup_month"] = df["tpep_pickup_datetime"].dt.month  
  
month_order = list(range(1, 13))  
  
plt.figure(figsize=(10, 5))  
df["pickup_month"].value_counts().reindex(month_order).plot(kind="bar")  
plt.xlabel("Month")  
plt.ylabel("Number of Pickups")  
plt.title("Distribution of Taxi Pickups by Month")  
plt.xticks(ticks=range(12), labels=["Jan", "Feb", "Mar", "Apr", "May",  
"Jun",  
"Jul", "Aug", "Sep", "Oct",  
"Nov", "Dec"], rotation=0)  
plt.grid(axis="y", linestyle="--", alpha=0.7)  
plt.show()
```



Financial Analysis

Take a look at the financial parameters like `fare_amount`, `tip_amount`, `total_amount`, and also `trip_distance`. Do these contain zero/negative values?

Analyse the above parameters

```
(df['fare_amount'] < 0.01).value_counts()
```

```
fare_amount
False      250879
True         28
Name: count, dtype: int64
```

```
df = df[df['fare_amount'] != 0]
(df['fare_amount'] < 0.01).value_counts()
```

```
fare_amount
False      250879
Name: count, dtype: int64
```

```
(df['total_amount'] < 0.01).value_counts()
```

```
total_amount
False      250879
Name: count, dtype: int64
```

```
(df['trip_distance'] < 0.01).value_counts()
```

```
trip_distance
False      248098
```

```

True      2781
Name: count, dtype: int64

df = df[df['trip_distance'] != 0]
(df['trip_distance'] < 0.01).value_counts()

trip_distance
False      248098
Name: count, dtype: int64

```

Do you think it is beneficial to create a copy DataFrame leaving out the zero values from these?

Yes removed the Zero values from fare_amount, total_amount, and also trip_distance (Top_Amount not removed as it is optional amount from the customer it can be zero too)

3.1.3 [2 marks] Filter out the zero values from the above columns.

Note: The distance might be 0 in cases where pickup and drop is in the same zone. Do you think it is suitable to drop such cases of zero distance?

```

df["PU_DO_Diff"] = df["PULocationID"] - df["DOLocationID"]

pu_do_diff_counts = df["PU_DO_Diff"].value_counts().reset_index()
pu_do_diff_counts.columns = ["PU_DO_Diff", "count"]

zero_diff_counts = pu_do_diff_counts[pu_do_diff_counts["PU_DO_Diff"]
== 0]

print(zero_diff_counts)

   PU_DO_Diff  count
0           0  11634

# Create a df with non zero entries for the selected parameters.
df = df[df["PU_DO_Diff"] != 0]

pu_do_diff_counts = df["PU_DO_Diff"].value_counts().reset_index()
pu_do_diff_counts.columns = ["PU_DO_Diff", "count"]

zero_diff_counts = pu_do_diff_counts[pu_do_diff_counts["PU_DO_Diff"]
== 0]
print(zero_diff_counts)

columns_to_remove = ["pickup_hour", "pickup_day", "pickup_month",
"PU_DO_Diff"]

```

```
Empty DataFrame
Columns: [PU_DO_Diff, count]
Index: []
```

3.1.4 [3 marks] Analyse the monthly revenue (`total_amount`) trend

```
# Group data by month and analyse monthly revenue
```


```
df['tpep_pickup_datetime'] =  
pd.to_datetime(df['tpep_pickup_datetime'])
```

```
df['Year-Month'] = df['tpep_pickup_datetime'].dt.to_period('M')
```

```
monthly_revenue = df.groupby('Year-Month')  
['total_amount'].sum().reset_index()
```

```
monthly_revenue['Year-Month'] = monthly_revenue['Year-  
Month'].astype(str)
```

```
plt.figure(figsize=(12, 6))  
plt.plot(monthly_revenue['Year-Month'],  
monthly_revenue['total_amount'], marker='o', linestyle='-')  
plt.xlabel("Month")  
plt.ylabel("Total Revenue ($)")  
plt.title("Monthly Revenue Trend")  
plt.xticks(rotation=45)  
plt.grid(True)  
plt.show()
```




3.1.5 [3 marks] Show the proportion of each quarter of the year in the revenue

Calculate proportion of each quarter

```
df["pickup_quarter"] =  
pd.to_datetime(df["tpep_pickup_datetime"]).dt.quarter  
  
quarterly_revenue = df.groupby("pickup_quarter")["total_amount"].sum()  
proportion = (quarterly_revenue / quarterly_revenue.sum()) * 100  
  
plt.figure(figsize=(8, 8))  
plt.pie(proportion, labels=["Q1", "Q2", "Q3", "Q4"], autopct='%1.1f%%', colors=["lightblue", "lightgreen", "salmon", "gold"])  
plt.title("Revenue Proportion by Quarter")  
plt.show()
```

Revenue Proportion by Quarter



3.1.6 [3 marks] Visualise the relationship between `trip_distance` and `fare_amount`. Also find the correlation value for these two.

Hint: You can leave out the trips with `trip_distance = 0`

```
df = df[df["fare_amount"] <= 700]
# Show how trip fare is affected by distance

df_filtered = df[df["trip_distance"] > 0]
```


```

correlation_value =
df_filtered["trip_distance"].corr(df_filtered["fare_amount"])
print(f"Correlation between trip distance and fare amount:
{correlation_value:.2f}")

plt.figure(figsize=(8, 5))
sns.scatterplot(x=df_filtered["trip_distance"],
y=df_filtered["fare_amount"], alpha=0.5)
plt.xlabel("Trip Distance (miles)")
plt.ylabel("Fare Amount ($)")
plt.title(f"Trip Distance vs. Fare Amount\nCorrelation:
{correlation_value:.2f}")
plt.grid(True)
plt.show()

```

Correlation between trip distance and fare amount: 0.96



3.1.7 [5 marks] Find and visualise the correlation between:

1. `fare_amount` and trip duration (pickup time to dropoff time)
2. `fare_amount` and `passenger_count`
3. `tip_amount` and `trip_distance`


```

# Show relationship between fare and trip duration
df["tpep_pickup_datetime"] =
pd.to_datetime(df["tpep_pickup_datetime"])
df["tpep_dropoff_datetime"] =
pd.to_datetime(df["tpep_dropoff_datetime"])

df["trip_duration"] = (df["tpep_dropoff_datetime"] -
df["tpep_pickup_datetime"]).dt.total_seconds() / 60


df_filtered = df[df["trip_duration"] > 0]

correlation_value =
df_filtered["fare_amount"].corr(df_filtered["trip_duration"])
print(f"Correlation between fare amount and trip duration:
{correlation_value:.2f}")


plt.figure(figsize=(8, 5))
sns.scatterplot(x=df_filtered["trip_duration"],
y=df_filtered["fare_amount"], alpha=0.5)
plt.xlabel("Trip Duration (minutes)")
plt.ylabel("Fare Amount ($)")
plt.title(f"Trip Duration vs. Fare Amount\nCorrelation:
{correlation_value:.2f}")
plt.grid(True)
plt.show()

```

Correlation between fare amount and trip duration: 0.28



```
# Show relationship between fare and number of passengers
plt.figure(figsize=(10, 6))
sns.boxplot(x=df_filtered["passenger_count"],
y=df_filtered["fare_amount"], color='red') # Corrected spelling
plt.xlabel("Passenger Count")
plt.ylabel("Fare Amount ($)")
plt.title("Fare Amount Distribution for Different Passenger Counts")
plt.grid(True)
plt.show()
```

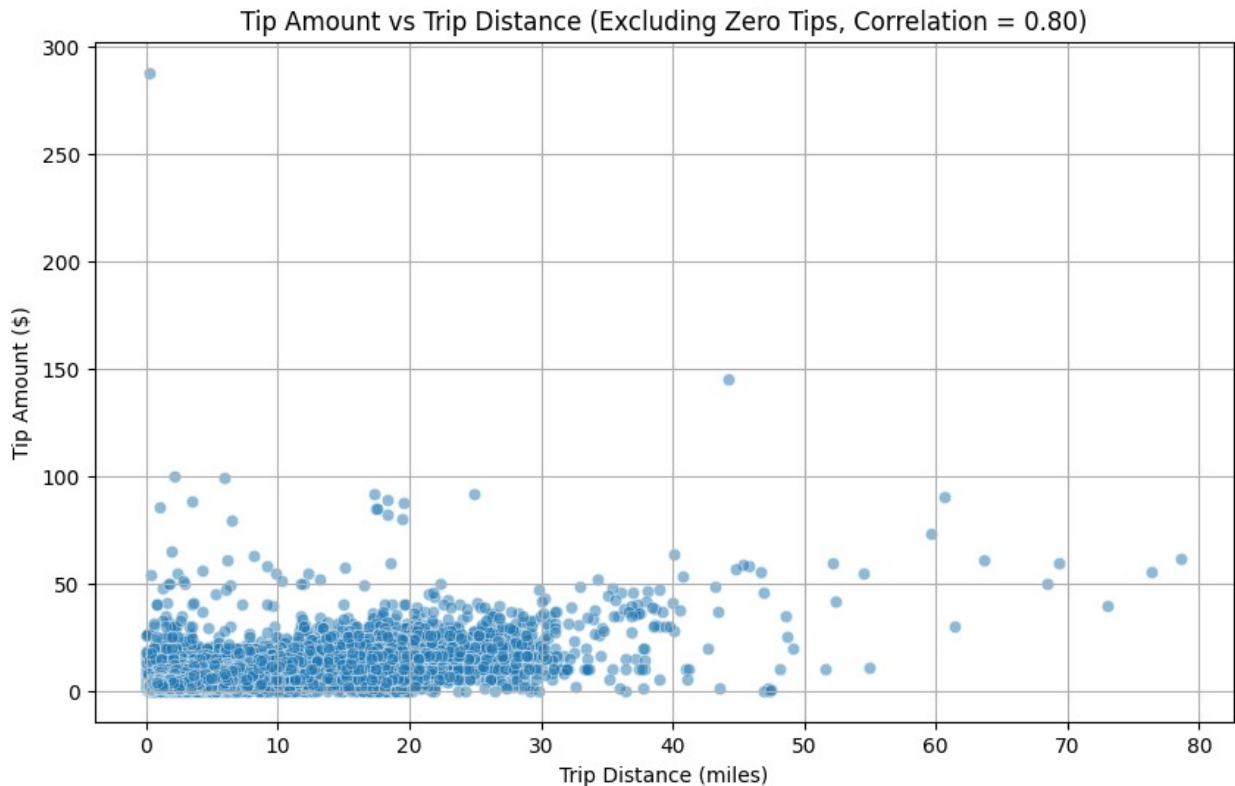


Show relationship between tip and trip distance

```
filtered_df = df[df['tip_amount'] > 0]
```

```
correlation_value =  
filtered_df['tip_amount'].corr(filtered_df['trip_distance'])
```

```
plt.figure(figsize=(10, 6))  
sns.scatterplot(x=filtered_df['trip_distance'],  
y=filtered_df['tip_amount'], alpha=0.5)  
plt.xlabel("Trip Distance (miles)")  
plt.ylabel("Tip Amount ($)")  
plt.title(f"Tip Amount vs Trip Distance (Excluding Zero Tips,  
Correlation = {correlation_value:.2f})")  
plt.grid(True)  
plt.show()
```




3.1.8 [3 marks] Analyse the distribution of different payment types (`payment_type`)

```
# Analyse the distribution of different payment types (payment_type).
payment_distribution = df['payment_type'].value_counts().reset_index()
payment_distribution.columns = ['Payment Type', 'Count']

plt.figure(figsize=(8, 6))
sns.barplot(x=payment_distribution['Payment Type'],
            y=payment_distribution['Count'], palette='viridis')

plt.xlabel("Payment Type")
plt.ylabel("Count")
plt.title("Distribution of Different Payment Types")
plt.grid(axis='y')
plt.show()
```



- 1= Credit card
- 2= Cash
- 3= No charge
- 4= Dispute

Geographical Analysis

For this, you have to use the *taxi_zones.shp* file from the *taxi_zones* folder.

There would be multiple files inside the folder (such as *.shx*, *.sbx*, *.sbn* etc). You do not need to import/read any of the files other than the shapefile, *taxi_zones.shp*.

Do not change any folder structure - all the files need to be present inside the folder for it to work.

The folder structure should look like this:

```
Taxi Zones
|- taxi_zones.shp.xml
|- taxi_zones.prj
|- taxi_zones.sbn
|- taxi_zones.shp
```

```
| - taxi_zones.dbf  
| - taxi_zones.shx  
| - taxi_zones.sbx
```

You only need to read the `taxi_zones.shp` file. The `shp` file will utilise the other files by itself.

We will use the *GeoPandas* library for geographical analysis

```
import geopandas as gpd
```

More about geopandas and shapefiles: [About](#)

Reading the shapefile is very similar to *Pandas*. Use `gpd.read_file()` function to load the data (`taxi_zones.shp`) as a `GeoDataFrame`. Documentation: [Reading and Writing Files](#)

```
!pip install geopandas
```

```
Requirement already satisfied: geopandas in c:\users\admin\anaconda3\lib\site-packages (1.0.1)  
Requirement already satisfied: numpy>=1.22 in c:\users\admin\anaconda3\lib\site-packages (from geopandas) (1.26.4)  
Requirement already satisfied: pyogrio>=0.7.2 in c:\users\admin\anaconda3\lib\site-packages (from geopandas) (0.10.0)  
Requirement already satisfied: packaging in c:\users\admin\anaconda3\lib\site-packages (from geopandas) (24.1)  
Requirement already satisfied: pandas>=1.4.0 in c:\users\admin\anaconda3\lib\site-packages (from geopandas) (2.2.2)  
Requirement already satisfied: pyproj>=3.3.0 in c:\users\admin\anaconda3\lib\site-packages (from geopandas) (3.7.1)  
Requirement already satisfied: shapely>=2.0.0 in c:\users\admin\anaconda3\lib\site-packages (from geopandas) (2.0.7)  
Requirement already satisfied: python-dateutil>=2.8.2 in c:\users\admin\anaconda3\lib\site-packages (from pandas>=1.4.0->geopandas) (2.9.0.post0)  
Requirement already satisfied: pytz>=2020.1 in c:\users\admin\anaconda3\lib\site-packages (from pandas>=1.4.0->geopandas) (2024.1)  
Requirement already satisfied: tzdata>=2022.7 in c:\users\admin\anaconda3\lib\site-packages (from pandas>=1.4.0->geopandas) (2023.3)  
Requirement already satisfied: certifi in c:\users\admin\anaconda3\lib\site-packages (from pyogrio>=0.7.2->geopandas) (2024.8.30)  
Requirement already satisfied: six>=1.5 in c:\users\admin\anaconda3\lib\site-packages (from python-dateutil>=2.8.2->pandas>=1.4.0->geopandas) (1.16.0)
```

3.1.9 [2 marks] Load the shapefile and display it.

```
# import geopandas as gpd  
import geopandas as gpd
```

```
shapefile_path = r"C:\Users\ADMIN\AL ML Coursesss\taxi_zones\taxi_zones.shp"
```

```
zones = gpd.read_file(shapefile_path)
```

```
zones.head()
```

	OBJECTID	Shape_Leng	Shape_Area	zone
0	1	0.116357	0.000782	Newark Airport
1	2	0.433470	0.004866	Jamaica Bay
2	3	0.084341	0.000314	Allerton/Pelham Gardens
3	4	0.043567	0.000112	Alphabet City
4	5	0.092146	0.000498	Arden Heights

	borough	geometry
0	EWB	POLYGON ((933100.918 192536.086, 933091.011 19...
1	Queens	MULTIPOLYGON (((1033269.244 172126.008, 103343...
2	Bronx	POLYGON ((1026308.77 256767.698, 1026495.593 2...
3	Manhattan	POLYGON ((992073.467 203714.076, 992068.667 20...
4	Staten Island	POLYGON ((935843.31 144283.336, 936046.565 144...

Now, if you look at the DataFrame created, you will see columns like: `OBJECTID`, `Shape_Leng`, `Shape_Area`, `zone`, `LocationID`, `borough`, `geometry`.

Now, the `locationID` here is also what we are using to mark pickup and drop zones in the trip records.

The geometric parameters like shape length, shape area and geometry are used to plot the zones on a map.

This can be easily done using the `plot()` method.

```
print(zones.info())
```

```
zones.plot()
```

```
<class 'geopandas.geodataframe.GeoDataFrame'>
```

```
RangeIndex: 263 entries, 0 to 262
```


```
Data columns (total 7 columns):
```

#	Column	Non-Null Count	Dtype
0	OBJECTID	263 non-null	int32
1	Shape_Leng	263 non-null	float64
2	Shape_Area	263 non-null	float64
3	zone	263 non-null	object

```

4  LocationID  263 non-null  int32
5  borough    263 non-null  object
6  geometry   263 non-null  geometry
dtypes: float64(2), geometry(1), int32(2), object(2)
memory usage: 12.5+ KB
None
<Axes: >

```



Now, you have to merge the trip records and zones data using the location IDs.

3.1.10 [3 marks] Merge the zones data into trip data using the `locationID` and `PULocationID` columns.

```

df.info()

<class 'pandas.core.frame.DataFrame'>
Index: 236463 entries, 0 to 265499
Data columns (total 26 columns):
#   Column                                Non-Null Count  Dtype  
---  -
0   VendorID                             236463 non-null int64   
1   tpep_pickup_datetime                 236463 non-null datetime64[ns]
2   tpep_dropoff_datetime                236463 non-null datetime64[ns]

```



```

3  passenger_count      236463 non-null float64
4  trip_distance        236463 non-null float64
5  RatecodeID           236463 non-null float64
6  store_and_fwd_flag   236463 non-null object
7  PULocationID         236463 non-null int64
8  DOLocationID         236463 non-null int64
9  payment_type         236463 non-null int64
10 fare_amount          236463 non-null float64
11 extra                236463 non-null float64
12 mta_tax              236463 non-null float64
13 tip_amount           236463 non-null float64
14 tolls_amount         236463 non-null float64
15 improvement_surcharge 236463 non-null float64
16 total_amount         236463 non-null float64
17 congestion_surcharge 236463 non-null float64
18 Combined_Airport_Fee 236463 non-null float64
19 pickup_hour          236463 non-null int32
20 pickup_day           236463 non-null int32
21 pickup_month         236463 non-null int32
22 PU_DO_Diff           236463 non-null int64
23 Year-Month           236463 non-null period[M]
24 pickup_quarter       236463 non-null int32
25 trip_duration        236463 non-null float64
dtypes: datetime64[ns](2), float64(13), int32(4), int64(5), object(1),
period[M](1)
memory usage: 45.1+ MB

```

#") Merge zones and trip records using locationID and PULocationID

```

zones["LocationID"] = zones["LocationID"].astype(int)
df["PULocationID"] = df["PULocationID"].astype(int)

merged_df = df.merge(zones, left_on="PULocationID",
right_on="LocationID", how="left")

print(merged_df)

```

	VendorID	tpep_pickup_datetime	tpep_dropoff_datetime
passenger_count \			
0	1	2023-01-05 07:50:08	2023-01-05 08:02:04
2.0			
1	2	2023-01-17 07:47:24	2023-01-17 08:00:50
5.0			
2	2	2023-01-25 21:57:59	2023-01-25 22:00:33
1.0			
3	2	2023-01-09 19:36:54	2023-01-09 19:52:01
2.0			
4	1	2023-01-11 22:19:13	2023-01-11 22:32:37
1.0			
...

```

...
236465      1  2023-09-25 23:11:46  2023-09-25 23:19:38
1.0
236466      2  2023-09-11 07:42:12  2023-09-11 07:49:04
1.0
236467      2  2023-09-20 19:10:57  2023-09-20 19:38:54
1.0
236468      1  2023-09-30 13:48:26  2023-09-30 14:03:59
1.0
236469      2  2023-09-11 18:13:38  2023-09-11 18:27:04
1.0

```

```

      trip_distance  RatecodeID  store_and_fwd_flag  PULocationID  \
0                1.90          1.0                N          239
1                1.86          1.0                N          239
2                0.50          1.0                N          162
3                2.56          1.0                N          162
4                2.80          1.0                N          164
...              ...          ...              ...          ...
236465          1.10          1.0                N          140
236466          1.15          1.0                N          141
236467          9.49          1.0                N          234
236468          1.20          1.0                N          113
236469          0.77          1.0                N          107

```

```

      DOLocationID  payment_type  ...  Year-Month  pickup_quarter  \
0                236            1  ...    2023-01            1
1                162            1  ...    2023-01            1
2                170            1  ...    2023-01            1
3                262            1  ...    2023-01            1
4                231            1  ...    2023-01            1
...              ...          ...    ...          ...
236465          263            2  ...    2023-09            3
236466          262            1  ...    2023-09            3
236467          138            1  ...    2023-09            3
236468          211            1  ...    2023-09            3
236469          186            1  ...    2023-09            3

```

```

      trip_duration  OBJECTID  Shape_Leng  Shape_Area  \
0          11.933333      239.0    0.063626    0.000205
1          13.433333      239.0    0.063626    0.000205
2           2.566667      162.0    0.035270    0.000048
3          15.116667      162.0    0.035270    0.000048
4          13.400000      164.0    0.035772    0.000056
...              ...          ...    ...          ...
236465          7.866667      140.0    0.047584    0.000114
236466          6.866667      141.0    0.041514    0.000077
236467         27.950000      234.0    0.036072    0.000073
236468         15.550000      113.0    0.032745    0.000058
236469         13.433333      107.0    0.038041    0.000075

```

	zone	LocationID	borough	\
0	Upper West Side South	239.0	Manhattan	
1	Upper West Side South	239.0	Manhattan	
2	Midtown East	162.0	Manhattan	
3	Midtown East	162.0	Manhattan	
4	Midtown South	164.0	Manhattan	
...	
236465	Lenox Hill East	140.0	Manhattan	
236466	Lenox Hill West	141.0	Manhattan	
236467	Union Sq	234.0	Manhattan	
236468	Greenwich Village North	113.0	Manhattan	
236469	Gramercy	107.0	Manhattan	

	geometry
0	POLYGON ((991168.979 226252.992, 991955.565 22...
1	POLYGON ((991168.979 226252.992, 991955.565 22...
2	POLYGON ((992224.354 214415.293, 992096.999 21...
3	POLYGON ((992224.354 214415.293, 992096.999 21...
4	POLYGON ((988787.425 210315.593, 988662.868 21...
...	...
236465	POLYGON ((995735.062 215619.835, 995670.105 21...
236466	POLYGON ((994839.073 216123.698, 994786.74 216...
236467	POLYGON ((987029.847 207022.299, 987048.27 206...
236468	POLYGON ((986643.64 204346.324, 986592.535 204...
236469	POLYGON ((989131.643 205749.904, 989084.531 20...

[236470 rows x 33 columns]

3.1.11 [3 marks] Group data by location IDs to find the total number of trips per location ID

```
# Group data by location and calculate the number of trips

trip_counts_by_location =
merged_df.groupby("zone").size().reset_index(name="Number of Trips")

print(trip_counts_by_location)
```

	zone	Number of Trips
0	Alphabet City	237
1	Astoria	80
2	Baisley Park	71
3	Bath Beach	2
4	Battery Park	118
...
183	Williamsburg (South Side)	39
184	Woodside	31
185	World Trade Center	1249
186	Yorkville East	3137
187	Yorkville West	4543

```
[188 rows x 2 columns]
```

3.1.12 [2 marks] Now, use the grouped data to add number of trips to the GeoDataFrame.

We will use this to plot a map of zones showing total trips per zone.

```
# Merge trip counts back to the zones GeoDataFrame
df['trip_duration']=(df['tpep_dropoff_datetime'] -
df['tpep_pickup_datetime']).dt.total_seconds()/60

df_filtered = df[df["trip_duration"] > 0]

correlation_fare_duration = df['fare_amount'].corr(df['trip_duration'])
correlation_fare_passenger=df
['fare_amount'].corr(df['passenger_count'])

print(f"correlation_fare_duration:{correlation_fare_duration}")
print(f"correlation_fare_passenger:{correlation_fare_passenger}")

# Function to safely convert WKT strings to Shapely geometries
def convert_geometry(geom):
    if isinstance(geom, str): # Convert only if it's a string (WKT
format)
        return wkt.loads(geom)
    return geom # Keep existing Polygon/MultiPolygon objects
unchanged

zones['geometry'] = zones['geometry'].apply(convert_geometry)

geometry_types = zones['geometry'].apply(lambda x: x.geom_type if x
else "Invalid").value_counts()
geometry_types

geometry
Polygon          240
MultiPolygon      23
Name: count, dtype: int64

from shapely.geometry import Polygon, MultiPolygon

def convert_to_polygon(geom):
    if isinstance(geom, MultiPolygon):
        return max(geom.geoms, key=lambda g: g.area) # Keep the
largest polygon
    return geom # Keep existing Polygons unchanged

zones['geometry'] = zones['geometry'].apply(convert_to_polygon)
```

```

zones = zones.dropna(subset=['geometry'])

gdf = gpd.GeoDataFrame(zones, geometry='geometry', crs="EPSG:4326")

geometry_types_after = gdf['geometry'].apply(lambda x: x.geom_type if
x else "Invalid").value_counts()
geometry_types_after

geometry
Polygon      263
Name: count, dtype: int64

import geopandas as gpd
gdf = gpd.GeoDataFrame(zones, geometry='geometry', crs="EPSG:4326")

gdf = gdf.merge(trip_counts_by_location, on="zone", how="left")

gdf["Number of Trips"] = gdf["Number of Trips"].fillna(0)

```

The next step is creating a color map (choropleth map) showing zones by the number of trips taken.

Again, you can use the `zones.plot()` method for this. [Plot Method GPD](#)

But first, you need to define the figure and axis for the plot.

```
fig, ax = plt.subplots(1, 1, figsize = (12, 10))
```

This function creates a figure (fig) and a single subplot (ax)

After setting up the figure and axis, we can proceed to plot the GeoDataFrame on this axis. This is done in the next step where we use the plot method of the GeoDataFrame.

You can define the following parameters in the `zones.plot()` method:

```

column = '',
ax = ax,
legend = True,
legend_kwds = {'label': "label", 'orientation':
"<horizontal/vertical>"}

```

To display the plot, use `plt.show()`.

3.1.13 [3 marks] Plot a color-coded map showing zone-wise trips

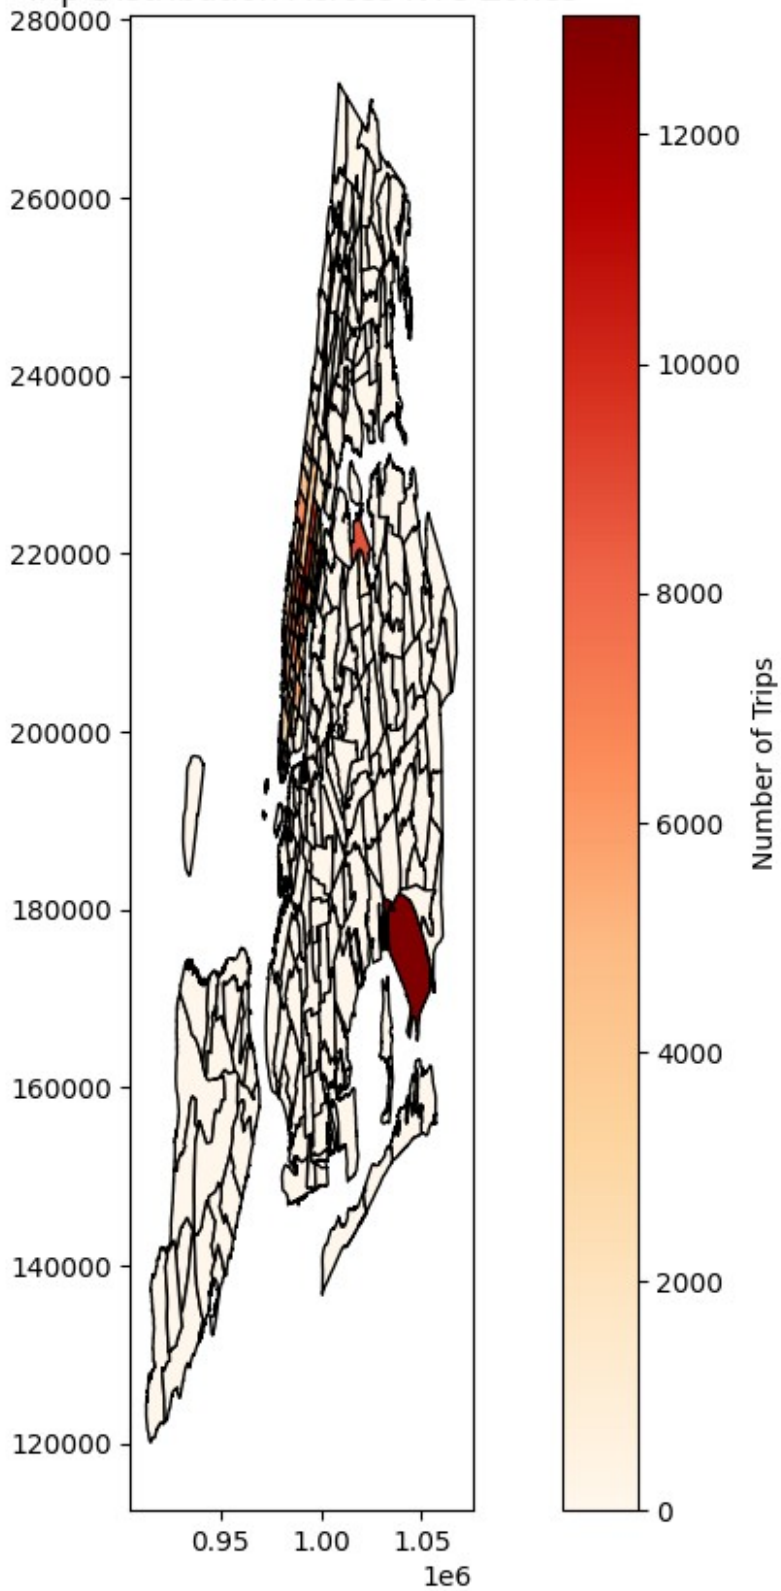
```
# Define figure and axis

fig, ax = plt.subplots(1, 1, figsize=(12, 10))

gdf.plot(column="Number of Trips", ax=ax, legend=True, cmap="OrRd",
          legend_kwds={'label': "Number of Trips", 'orientation':
"vertical"},
          edgecolor="black")

ax.set_title("Trip Distribution Across NYC Zones")
plt.show()
```

Trip Distribution Across NYC Zones



```
# can you try displaying the zones DF sorted by the number of trips?
```

Here we have completed the temporal, financial and geographical analysis on the trip records.

Compile your findings from general analysis below:

You can consider the following points:

- Busiest hours, days and months
- Trends in revenue collected
- Trends in quarterly revenue
- How fare depends on trip distance, trip duration and passenger counts
- How tip amount depends on trip distance
- Busiest zones

3.2 Detailed EDA: Insights and Strategies

[50 marks]

Having performed basic analyses for finding trends and patterns, we will now move on to some detailed analysis focussed on operational efficiency, pricing strategies, and customer experience.

Operational Efficiency

Analyze variations by time of day and location to identify bottlenecks or inefficiencies in routes

3.2.1 [3 marks] Identify slow routes by calculating the average time taken by cabs to get from one zone to another at different hours of the day.

Speed on a route X for hour $Y = (\text{distance of the route } X / \text{average trip duration for hour } Y)$

```
# Find routes which have the slowest speeds at different times of the day
```

How does identifying high-traffic, high-demand routes help us?

3.2.2 [3 marks] Calculate the number of trips at each hour of the day and visualise them. Find the busiest hour and show the number of trips for that hour.

```
# Visualise the number of trips per hour and find the busiest hour
```

```
df["tpep_pickup_datetime"] =  
pd.to_datetime(df["tpep_pickup_datetime"])
```

```
df["hour_of_day"] = df["tpep_pickup_datetime"].dt.hour
```




```
trips_per_hour =
df.groupby("hour_of_day").size().reset_index(name="Total_Trips")

busiest_hour =
trips_per_hour.loc[trips_per_hour["Total_Trips"].idxmax()]

plt.figure(figsize=(12, 6))
plt.bar(trips_per_hour["hour_of_day"], trips_per_hour["Total_Trips"],
alpha=0.75)
plt.xlabel("Hour of the Day")
plt.ylabel("Number of Trips")
plt.title("Number of Taxi Trips at Each Hour of the Day")
plt.xticks(range(24))
plt.grid(axis="y", linestyle="--", alpha=0.7)

busiest_hour_value = busiest_hour["hour_of_day"]
busiest_hour_trips = busiest_hour["Total_Trips"]
plt.axvline(busiest_hour_value, color="red", linestyle="--",
label=f"Busiest Hour: {busiest_hour_value} ({busiest_hour_trips}
trips)")
plt.legend()

plt.show()
```



Remember, we took a fraction of trips. To find the actual number, you have to scale the number up by the sampling ratio.

3.2.3 [2 mark] Find the actual number of trips in the five busiest hours

```
# Scale up the number of trips

# Fill in the value of your sampling fraction and use that to scale up
the numbers
sample_fraction =0.1
```

3.2.4 [3 marks] Compare hourly traffic pattern on weekdays. Also compare for weekend.

```
# Compare traffic trends for the week days and weekends

df["day_of_week"] = df["tpep_pickup_datetime"].dt.dayofweek

df["day_type"] = df["day_of_week"].apply(lambda x: "Weekend" if x >= 5
else "Weekday")

traffic_trends = df.groupby(["day_type",
"hour_of_day"]).size().reset_index(name="Total_Trips")


plt.figure(figsize=(12, 6))

weekday_trends = traffic_trends[traffic_trends["day_type"] ==
"Weekday"]
plt.plot(weekday_trends["hour_of_day"], weekday_trends["Total_Trips"],
marker='o', label="Weekdays")

weekend_trends = traffic_trends[traffic_trends["day_type"] ==
"Weekend"]
plt.plot(weekend_trends["hour_of_day"], weekend_trends["Total_Trips"],
marker='o', linestyle="--", label="Weekends")

plt.xlabel("Hour of the Day")
plt.ylabel("Number of Trips")
plt.title("Comparison of Traffic Trends: Weekdays vs Weekends")
plt.xticks(range(24))
plt.legend()
plt.grid(axis="y", linestyle="--", alpha=0.7)

plt.show()
```



What can you infer from the above patterns? How will finding busy and quiet hours for each day help us?

3.2.5 [3 marks] Identify top 10 zones with high hourly pickups. Do the same for hourly dropoffs. Show pickup and dropoff trends in these zones.

```
# Find top 10 pickup and dropoff zones
```

```
pickup_counts = df['PULocationID'].value_counts().reset_index()
pickup_counts.columns = ['LocationID', 'Pickup_Count']
```

```
dropoff_counts = df['DOLocationID'].value_counts().reset_index()
dropoff_counts.columns = ['LocationID', 'Dropoff_Count']
```

```
location_ratios = pd.merge(pickup_counts, dropoff_counts,
on='LocationID', how='outer').fillna(0)
```

```
location_ratios['Pickup_Dropoff_Ratio'] =
location_ratios['Pickup_Count'] / (location_ratios['Dropoff_Count'] +
1e-9)
```

```
top_10_ratios = location_ratios.nlargest(10, 'Pickup_Dropoff_Ratio')
bottom_10_ratios = location_ratios.nsmallest(10,
'Pickup_Dropoff_Ratio')
```

```
top_10_ratios, bottom_10_ratios
```

	LocationID	Pickup_Count	Dropoff_Count	Pickup_Dropoff_Ratio
(
56	59	1.0	0.0	1.000000e+09
67	70	1173.0	126.0	9.309524e+00
124	132	13039.0	2422.0	5.383567e+00
130	138	8694.0	2928.0	2.969262e+00
178	186	8634.0	5505.0	1.568392e+00
40	43	4107.0	2868.0	1.432008e+00
106	114	3270.0	2382.0	1.372796e+00
240	249	5394.0	3943.0	1.367994e+00
154	162	8748.0	6819.0	1.282886e+00
153	161	11366.0	9455.0	1.202115e+00,
	LocationID	Pickup_Count	Dropoff_Count	Pickup_Dropoff_Ratio
1	3	0.0	19.0	0.0
3	6	0.0	7.0	0.0
5	8	0.0	6.0	0.0
6	9	0.0	26.0	0.0
12	15	0.0	24.0	0.0
16	19	0.0	19.0	0.0
17	20	0.0	24.0	0.0
18	21	0.0	24.0	0.0
19	22	0.0	38.0	0.0
20	23	0.0	3.0	0.0)

3.2.6 [3 marks] Find the ratio of pickups and dropoffs in each zone. Display the 10 highest (pickup/drop) and 10 lowest (pickup/drop) ratios.

```
# Find the top 10 and bottom 10 pickup/dropoff ratios

pickup_dropoff_counts = df.groupby(['PULocationID',
'DOLocationID']).size().reset_index(name='Count')

pickup_counts = df['PULocationID'].value_counts().reset_index()
pickup_counts.columns = ['PULocationID', 'Total_Pickups']

dropoff_counts = df['DOLocationID'].value_counts().reset_index()
dropoff_counts.columns = ['DOLocationID', 'Total_Dropoffs']

pickup_dropoff_ratios = pickup_dropoff_counts.merge(pickup_counts,
on='PULocationID')
pickup_dropoff_ratios = pickup_dropoff_ratios.merge(dropoff_counts,
on='DOLocationID')

pickup_dropoff_ratios['Pickup_Dropoff_Ratio'] =
pickup_dropoff_ratios['Total_Pickups'] /
(pickup_dropoff_ratios['Total_Dropoffs'] + 1e-9)
```

```

top_10_ratios = pickup_dropoff_ratios.nlargest(10,
'Pickup_Dropoff_Ratio')
bottom_10_ratios = pickup_dropoff_ratios.nsmallest(10,
'Pickup_Dropoff_Ratio')

```

```

top_10_ratios, bottom_10_ratios

```

	PULocationID	DOLocationID	Count	Total_Pickups
Total_Dropoffs \				
2977	132	44	1	13039
1				
3104	132	187	1	13039
1				
3120	132	204	1	13039
1				
3539	138	253	1	8694
1				
2990	132	57	1	13039
2				
3093	132	176	1	13039
2				
3101	132	184	1	13039
2				
3366	138	57	1	8694
2				
3414	138	115	1	8694
2				
2956	132	23	3	13039
3				

	Pickup_Dropoff_Ratio			
2977	13038.999987			
3104	13038.999987			
3120	13038.999987			
3539	8693.999991			
2990	6519.499997			
3093	6519.499997			
3101	6519.499997			
3366	4346.999998			
3414	4346.999998			
2956	4346.333332			
	,			
	PULocationID	DOLocationID	Count	Total_Pickups
Total_Dropoffs	\			
400	31	237	1	1
9462				
1059	62	234	1	1
5760				
6044	228	186	1	1

5505				
461	38	161	1	2
9455				
5642	200	230	1	2
7547				
5643	200	239	1	2
6579				
5351	177	141	1	2
6189				
5196	167	75	1	1
2705				
5651	207	164	1	2
5298				
446	34	229	1	2
5179				

	Pickup_Dropoff_Ratio
400	0.000106
1059	0.000174
6044	0.000182
461	0.000212
5642	0.000265
5643	0.000304
5351	0.000323
5196	0.000370
5651	0.000378
446	0.000386)

3.2.7 [3 marks] Identify zones with high pickup and dropoff traffic during night hours (11PM to 5AM)

```
# During night hours (11pm to 5am) find the top 10 pickup and dropoff zones
# Note that the top zones should be of night hours and not the overall top zones
```

```
df['tpep_pickup_datetime'] =
pd.to_datetime(df['tpep_pickup_datetime'])
```

```
df['pickup_hour'] = df['tpep_pickup_datetime'].dt.hour
```

```
night_hours_data = df[(df['pickup_hour'] >= 23) | (df['pickup_hour']
<= 5)]
```

```
top_night_pickup_zones =
night_hours_data['PULocationID'].value_counts().head(10).reset_index()
top_night_pickup_zones.columns = ['PULocationID', 'Pickup_Count']
```

```

top_night_dropoff_zones =
night_hours_data['DOLocationID'].value_counts().head(10).reset_index()
top_night_dropoff_zones.columns = ['DOLocationID', 'Dropoff_Count']

top_night_pickup_zones, top_night_dropoff_zones

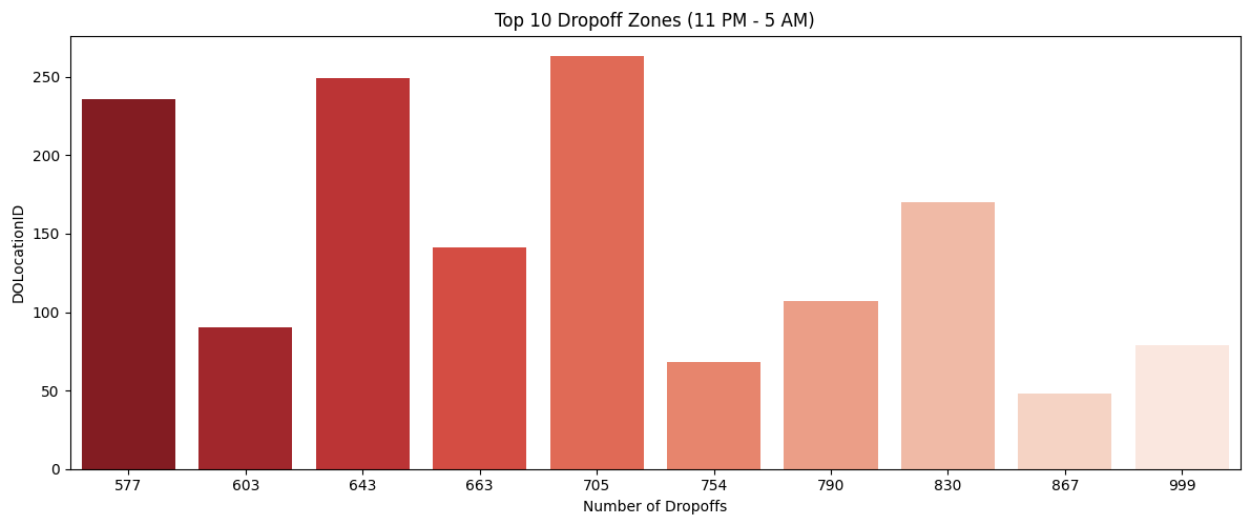
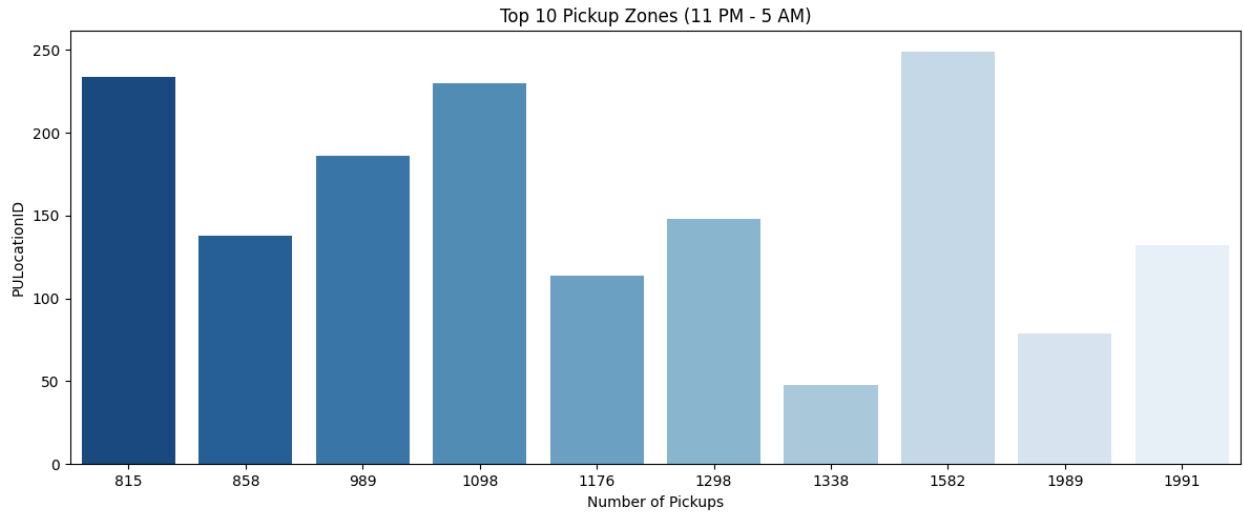
fig, axes = plt.subplots(2, 1, figsize=(12, 10))

sns.barplot(x='Pickup_Count', y='PULocationID',
data=top_night_pickup_zones, ax=axes[0], palette='Blues_r')
axes[0].set_title('Top 10 Pickup Zones (11 PM - 5 AM)')
axes[0].set_xlabel('Number of Pickups')
axes[0].set_ylabel('PULocationID')

sns.barplot(x='Dropoff_Count', y='DOLocationID',
data=top_night_dropoff_zones, ax=axes[1], palette='Reds_r')
axes[1].set_title('Top 10 Dropoff Zones (11 PM - 5 AM)')
axes[1].set_xlabel('Number of Dropoffs')
axes[1].set_ylabel('DOLocationID')

plt.tight_layout()
plt.show()

```



Now, let us find the revenue share for the night time hours and the day time hours. After this, we will move to deciding a pricing strategy.

3.2.8 [2 marks] Find the revenue share for nighttime and daytime hours.

```
# Filter for night hours (11 PM to 5 AM)
```

```
night_hours_data.head()
```

VendorID	tpep_pickup_datetime	tpep_dropoff_datetime	passenger_count
20	2023-01-04 01:07:43	2023-01-04 01:27:24	2
41	2023-01-28 00:11:13	2023-01-28 00:35:17	2
54	2023-01-25 23:04:10	2023-01-25 23:37:00	1
60	2023-01-22 04:14:54	2023-01-22 04:29:23	2

67 2 2023-01-30 01:30:07 2023-01-30 02:01:47
1.0

	trip_distance	RatecodeID	store_and_fwd_flag	PULocationID
DOLocationID \				
20	14.90	1.0	N	132
192				
41	5.44	1.0	N	50
232				
54	8.90	1.0	N	230
106				
60	3.19	1.0	N	232
25				
67	16.02	1.0	N	68
122				

	payment_type	...	pickup_hour	pickup_day	pickup_month	
PU_DOLocationID \						
20	1	...	1	2	1	-
60						
41	1	...	0	5	1	-
182						
54	1	...	23	2	1	
124						
60	2	...	4	6	1	
207						
67	2	...	1	0	1	-
54						

	Year-Month	pickup_quarter	trip_duration	hour_of_day
day_of_week \				
20	2023-01	1	19.683333	1
2				
41	2023-01	1	24.066667	0
5				
54	2023-01	1	32.833333	23
2				
60	2023-01	1	14.483333	4
6				
67	2023-01	1	31.666667	1
0				

	day_type
20	Weekday
41	Weekend
54	Weekday
60	Weekend
67	Weekday

[5 rows x 29 columns]

Pricing Strategy

3.2.9 [2 marks] For the different passenger counts, find the average fare per mile per passenger.

For instance, suppose the average fare per mile for trips with 3 passengers is 3 USD/mile, then the fare per mile per passenger will be 1 USD/mile.

Analyse the fare per mile per passenger for different passenger counts

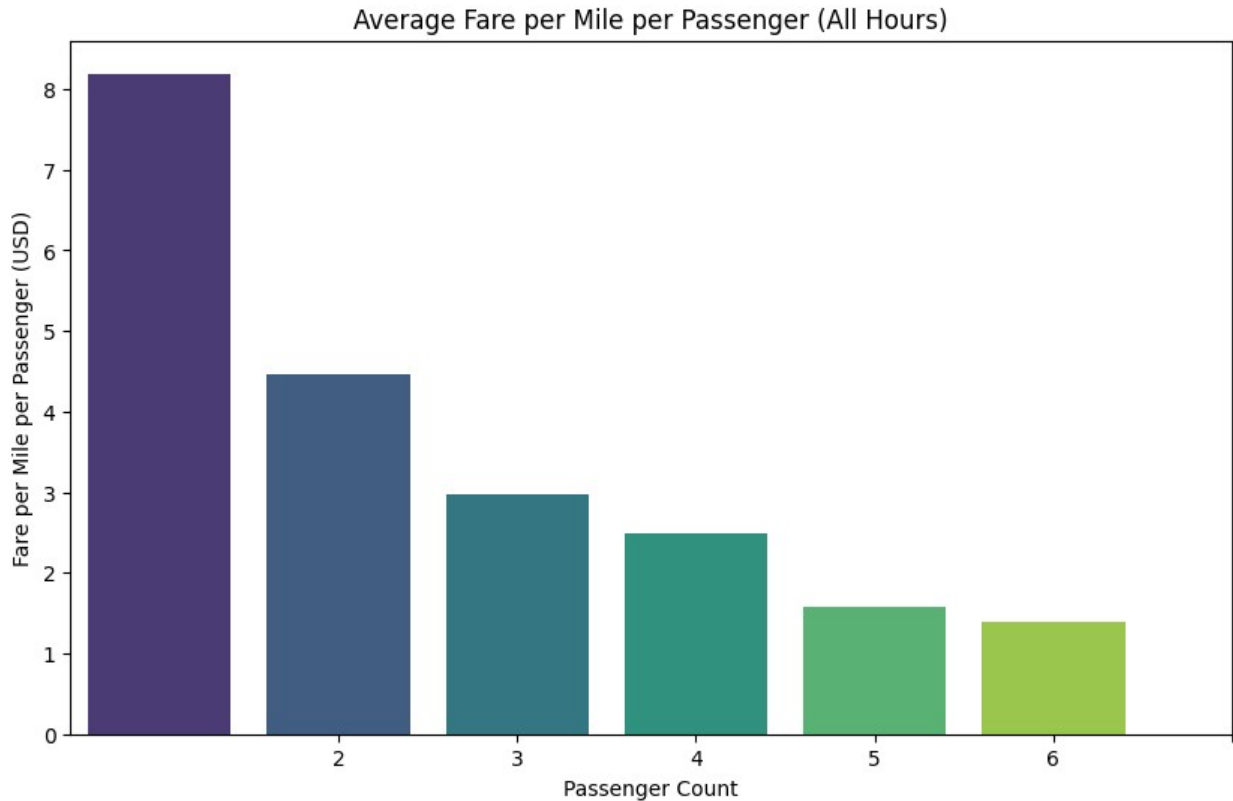
```
df = df[(df['trip_distance'] > 0) & (df['fare_amount'] > 0) &
(df['passenger_count'] > 0)]

df['fare_per_mile_per_passenger'] = df['fare_amount'] /
(df['trip_distance'] * df['passenger_count'])

avg_fare_per_mile_per_passenger['passenger_count'] =
avg_fare_per_mile_per_passenger['passenger_count'].astype(int)

plt.figure(figsize=(10, 6))
sns.barplot(x='passenger_count', y='fare_per_mile_per_passenger',
data=avg_fare_per_mile_per_passenger, palette='viridis')

plt.title('Average Fare per Mile per Passenger (All Hours)')
plt.xlabel('Passenger Count')
plt.ylabel('Fare per Mile per Passenger (USD)')
plt.xticks(ticks=range(1,
avg_fare_per_mile_per_passenger['passenger_count'].max() + 1))
plt.show()
```



3.2.10 [3 marks] Find the average fare per mile by hours of the day and by days of the week

Compare the average fare per mile for different days and for different times of the day

```
df['tpep_pickup_datetime'] =  
pd.to_datetime(df['tpep_pickup_datetime'])
```

```
df['day_of_week'] = df['tpep_pickup_datetime'].dt.day_name()
```

```
def categorize_time_of_day(hour):  
    if 5 <= hour < 12:  
        return 'Morning'  
    elif 12 <= hour < 17:  
        return 'Afternoon'  
    elif 17 <= hour < 21:  
        return 'Evening'  
    else:  
        return 'Night'
```

```
df['time_of_day'] =  
df['tpep_pickup_datetime'].dt.hour.apply(categorize_time_of_day)
```

```

valid_df = df[(df['trip_distance'] > 0) & (df['fare_amount'] > 0) &
(df['passenger_count'] > 0)]

valid_df['fare_per_mile'] = valid_df['fare_amount'] /
valid_df['trip_distance']

avg_fare_per_mile_by_day = valid_df.groupby('day_of_week')
['fare_per_mile'].mean().reindex(
    ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday',
'Saturday', 'Sunday']).reset_index()

avg_fare_per_mile_by_time = valid_df.groupby('time_of_day')
['fare_per_mile'].mean().reindex(
    ['Morning', 'Afternoon', 'Evening', 'Night']).reset_index()


fig, axes = plt.subplots(1, 2, figsize=(16, 6))

sns.barplot(x='day_of_week', y='fare_per_mile',
data=avg_fare_per_mile_by_day, palette='viridis', ax=axes[0])
axes[0].set_title('Average Fare per Mile by Day of the Week')
axes[0].set_xlabel('Day of the Week')
axes[0].set_ylabel('Fare per Mile (USD)')

sns.barplot(x='time_of_day', y='fare_per_mile',
data=avg_fare_per_mile_by_time, palette='plasma', ax=axes[1])
axes[1].set_title('Average Fare per Mile by Time of Day')
axes[1].set_xlabel('Time of Day')
axes[1].set_ylabel('Fare per Mile (USD)')

plt.tight_layout()
plt.show()

```



3.2.11 [3 marks] Analyse the average fare per mile for the different vendors for different hours of the day

Compare fare per mile for different vendors

```
df['tpep_pickup_datetime'] =
pd.to_datetime(df['tpep_pickup_datetime'])

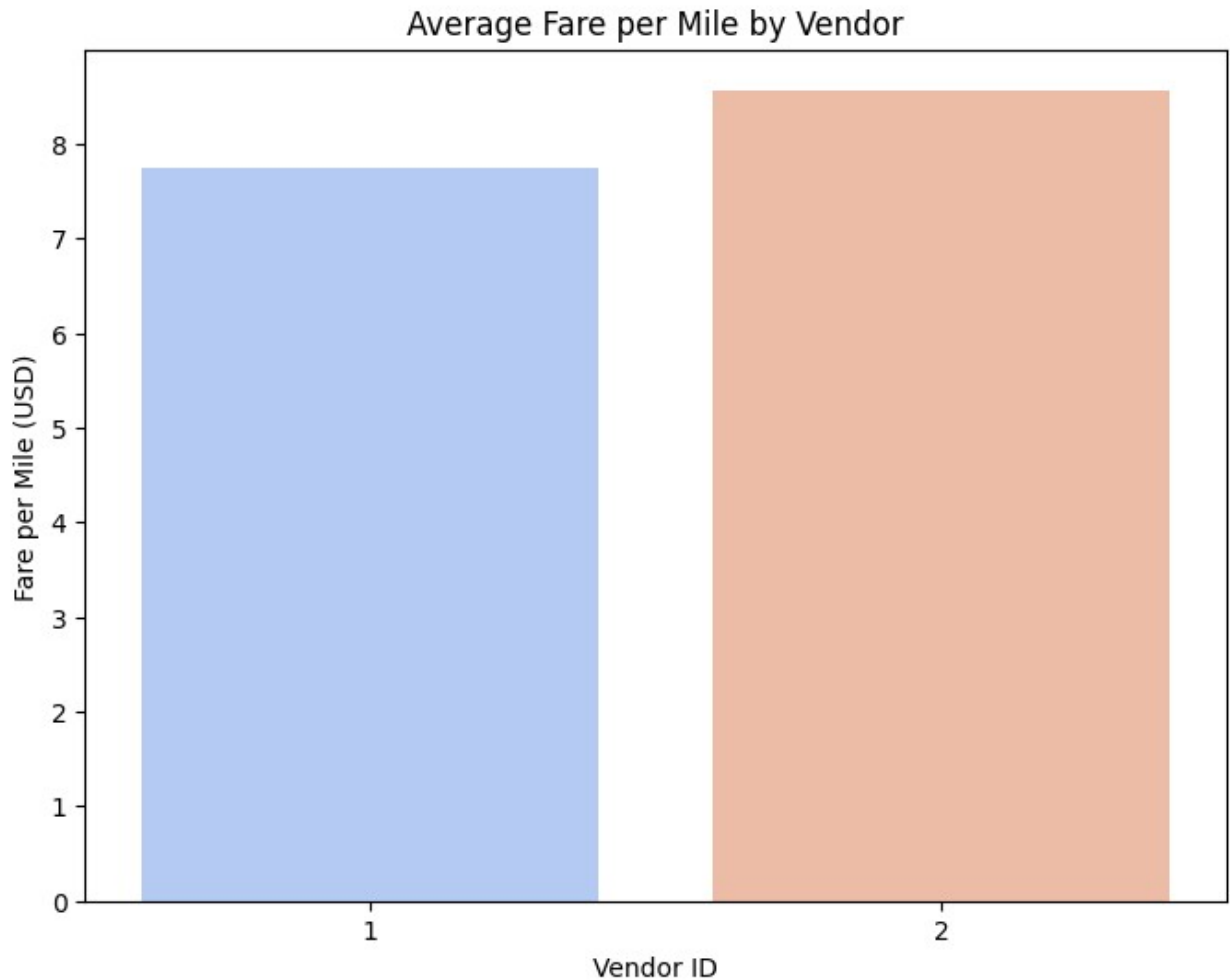
valid_df = df[(df['trip_distance'] > 0) & (df['fare_amount'] > 0)]

valid_df['fare_per_mile'] = valid_df['fare_amount'] /
valid_df['trip_distance']

avg_fare_per_mile_by_vendor = valid_df.groupby('VendorID')
['fare_per_mile'].mean().reset_index()

plt.figure(figsize=(8, 6))
sns.barplot(x='VendorID', y='fare_per_mile',
data=avg_fare_per_mile_by_vendor, palette='coolwarm')

plt.title('Average Fare per Mile by Vendor')
plt.xlabel('Vendor ID')
plt.ylabel('Fare per Mile (USD)')
plt.show()
```



3.2.12 [5 marks] Compare the fare rates of the different vendors in a tiered fashion. Analyse the average fare per mile for distances upto 2 miles. Analyse the fare per mile for distances from 2 to 5 miles. And then for distances more than 5 miles.

```
# Defining distance tiers
```

```
valid_df = df[(df['trip_distance'] > 0) & (df['fare_amount'] > 0)]
```

```
def distance_tier(distance):  
    if distance <= 2:  
        return 'Short (0-2 mi)'  
    elif distance <= 5:  
        return 'Medium (2-5 mi)'  
    else:  
        return 'Long (>5 mi)'
```

```

valid_df['distance_tier'] =
valid_df['trip_distance'].apply(distance_tier)

valid_df['fare_per_mile'] = valid_df['fare_amount'] /
valid_df['trip_distance']

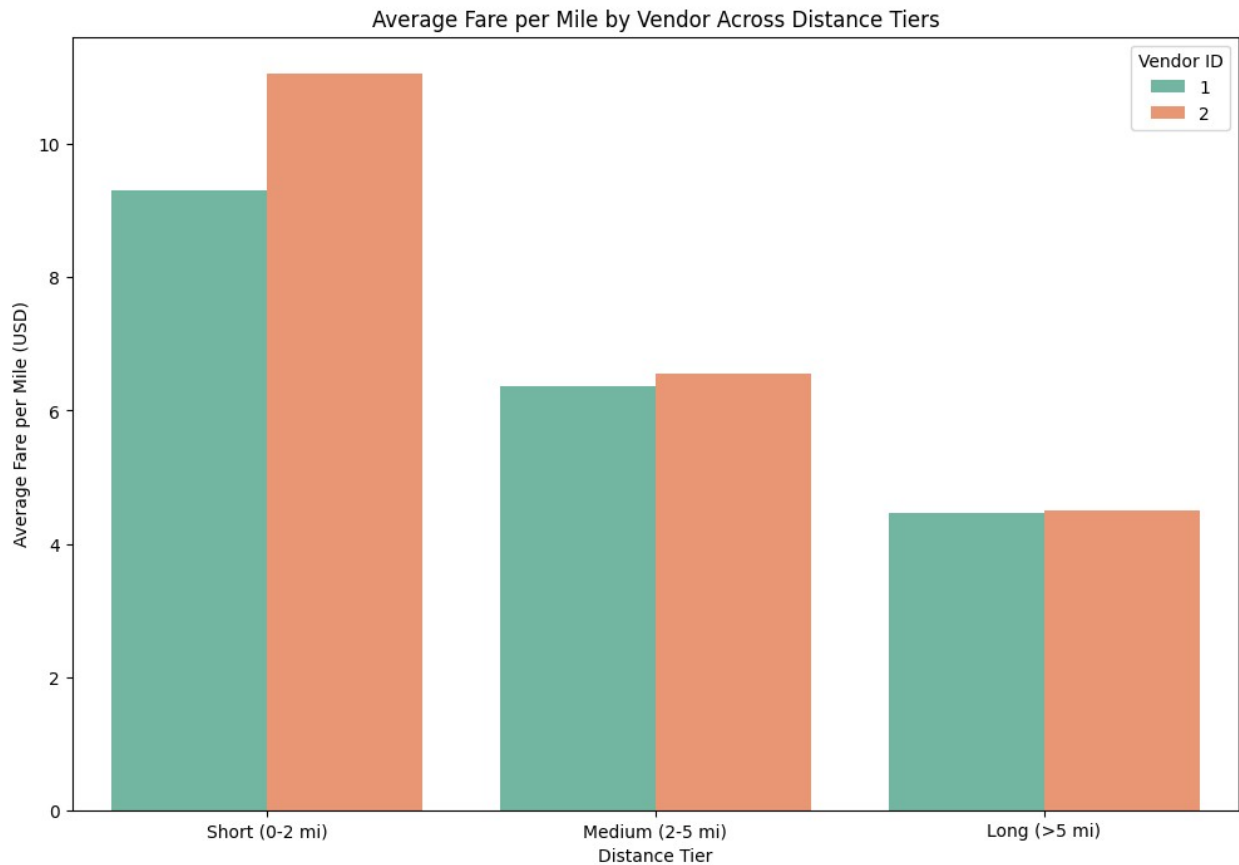
avg_fare_by_vendor_tier = valid_df.groupby(['VendorID',
'distance_tier'])['fare_per_mile'].mean().reset_index()

tier_order = ['Short (0-2 mi)', 'Medium (2-5 mi)', 'Long (>5 mi)']
avg_fare_by_vendor_tier['distance_tier'] =
pd.Categorical(avg_fare_by_vendor_tier['distance_tier'],
categories=tier_order, ordered=True)

plt.figure(figsize=(12, 8))
sns.barplot(x='distance_tier', y='fare_per_mile', hue='VendorID',
data=avg_fare_by_vendor_tier, palette='Set2')

plt.title('Average Fare per Mile by Vendor Across Distance Tiers')
plt.xlabel('Distance Tier')
plt.ylabel('Average Fare per Mile (USD)')
plt.legend(title='Vendor ID')
plt.show()

```



Customer Experience and Other Factors

3.2.13 [5 marks] Analyse average tip percentages based on trip distances, passenger counts and time of pickup. What factors lead to low tip percentages?

```
# Analyze tip percentages based on distances, passenger counts and pickup times
```

```
df['tpep_pickup_datetime'] =  
pd.to_datetime(df['tpep_pickup_datetime'])
```

```
valid_df = df[(df['fare_amount'] > 0) & (df['tip_amount'] >= 0)]
```

```
valid_df['tip_percentage'] = (valid_df['tip_amount'] /  
valid_df['fare_amount']) * 100
```

```
def distance_tier(distance):  
    if distance <= 2:  
        return 'Short (0-2 mi)'  
    elif distance <= 5:
```



```

        return 'Medium (2-5 mi)'
    else:
        return 'Long (>5 mi)'

valid_df['distance_tier'] =
valid_df['trip_distance'].apply(distance_tier)

def passenger_group(count):
    if count == 1:
        return '1 Passenger'
    elif count == 2:
        return '2 Passengers'
    elif 3 <= count <= 4:
        return '3-4 Passengers'
    else:
        return '5+ Passengers'

valid_df['passenger_group'] =
valid_df['passenger_count'].apply(passenger_group)

valid_df['pickup_hour'] = valid_df['tpep_pickup_datetime'].dt.hour

tip_by_distance = valid_df.groupby('distance_tier')
['tip_percentage'].mean().reset_index()

plt.figure(figsize=(8, 6))
sns.barplot(x='distance_tier', y='tip_percentage',
data=tip_by_distance, palette='Blues')
plt.title('Average Tip Percentage by Distance Tier')
plt.xlabel('Distance Tier')
plt.ylabel('Tip Percentage (%)')
plt.show()

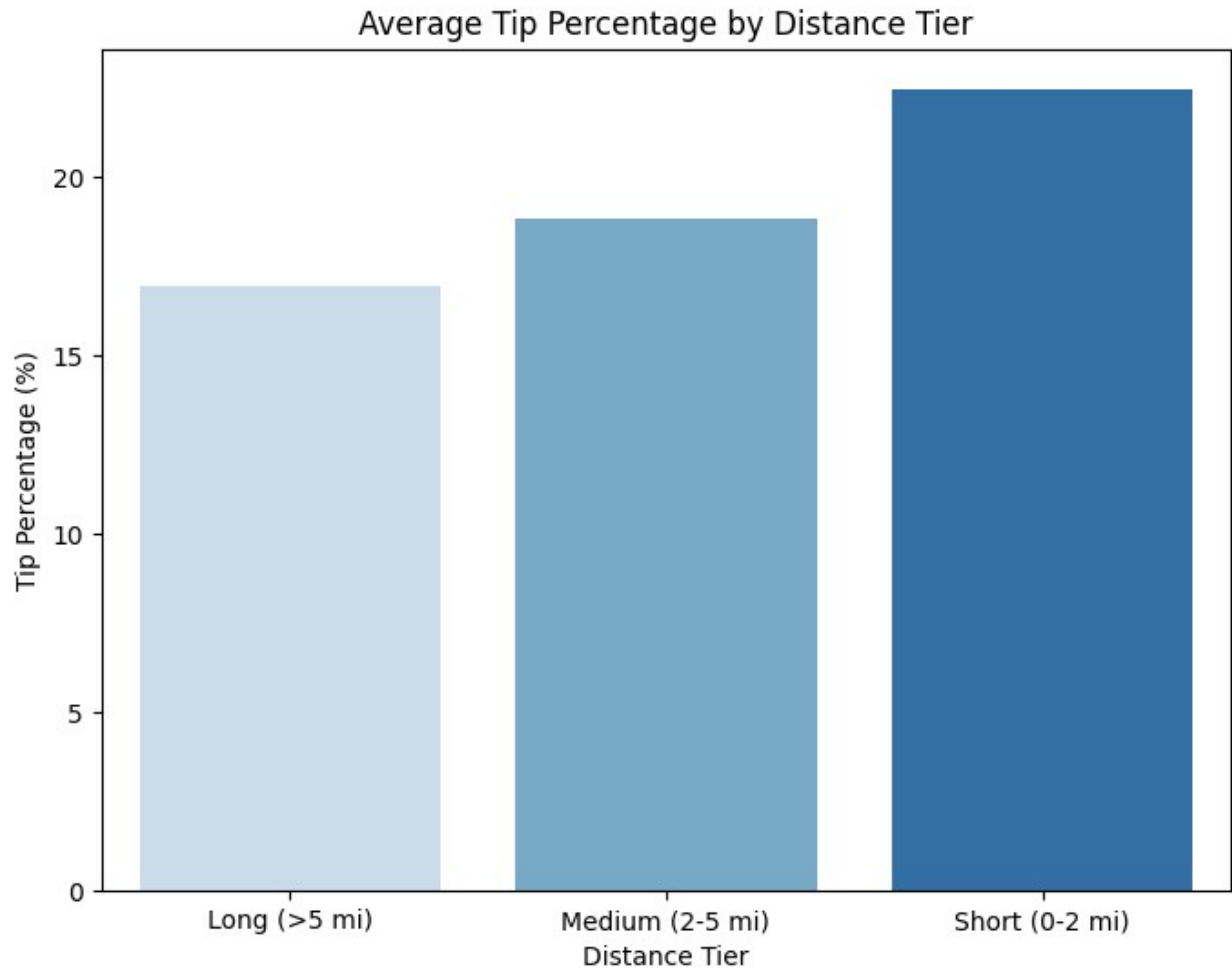
tip_by_passenger = valid_df.groupby('passenger_group')
['tip_percentage'].mean().reset_index()

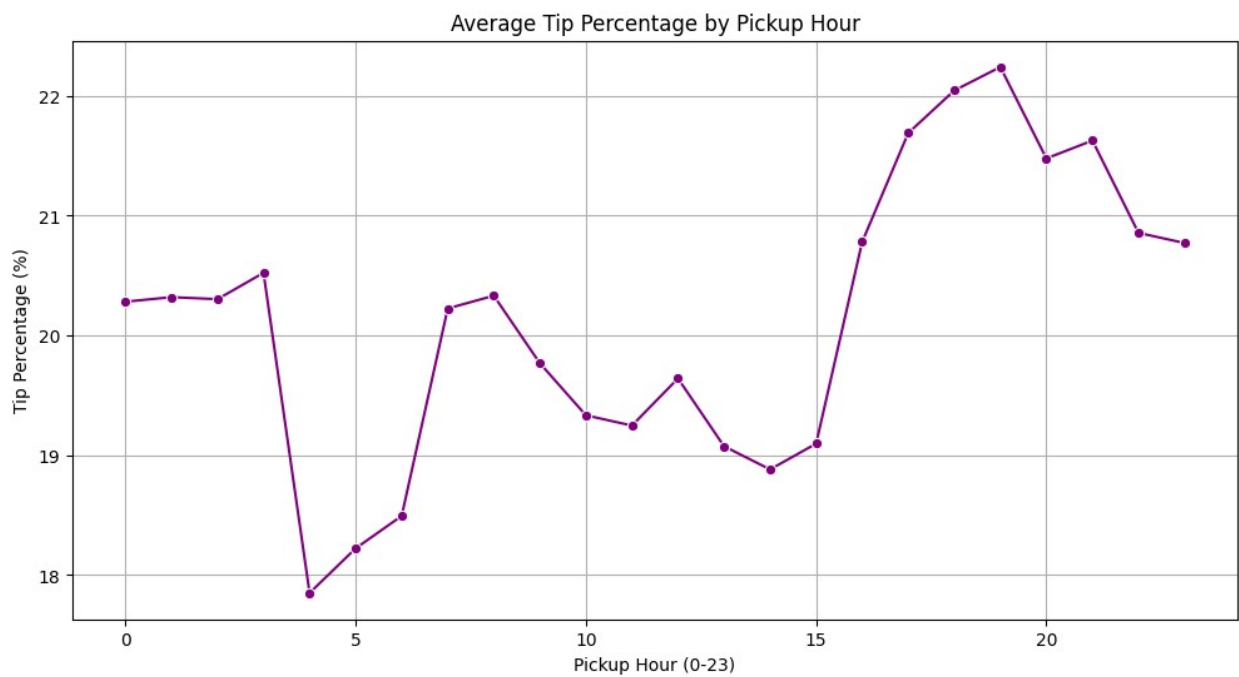
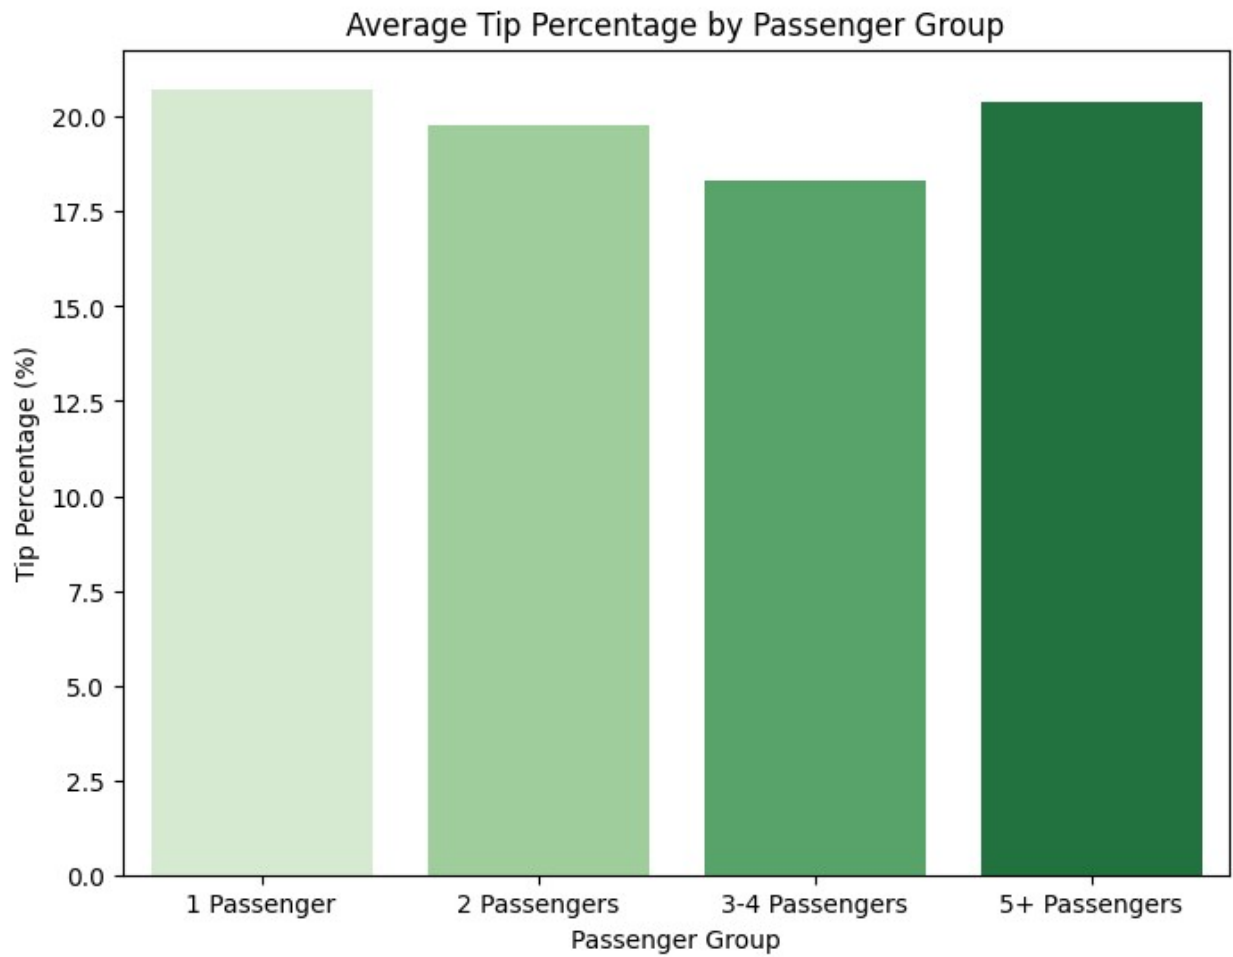
plt.figure(figsize=(8, 6))
sns.barplot(x='passenger_group', y='tip_percentage',
data=tip_by_passenger, palette='Greens')
plt.title('Average Tip Percentage by Passenger Group')
plt.xlabel('Passenger Group')
plt.ylabel('Tip Percentage (%)')
plt.show()

tip_by_hour = valid_df.groupby('pickup_hour')
['tip_percentage'].mean().reset_index()

```

```
plt.figure(figsize=(12, 6))
sns.lineplot(x='pickup_hour', y='tip_percentage', data=tip_by_hour,
marker='o', color='purple')
plt.title('Average Tip Percentage by Pickup Hour')
plt.xlabel('Pickup Hour (0-23)')
plt.ylabel('Tip Percentage (%)')
plt.grid(True)
plt.show()
```





Additional analysis [optional]: Let's try comparing cases of low tips with cases of high tips to find out if we find a clear aspect that drives up the tipping behaviours

```
# Compare trips with tip percentage < 10% to trips with tip percentage > 25%

df['tpep_pickup_datetime'] =
pd.to_datetime(df['tpep_pickup_datetime'])

valid_df = df[(df['fare_amount'] > 0) & (df['tip_amount'] >= 0)]

valid_df['tip_percentage'] = (valid_df['tip_amount'] /
valid_df['fare_amount']) * 100

low_tip_df = valid_df[valid_df['tip_percentage'] < 10]
high_tip_df = valid_df[valid_df['tip_percentage'] > 25]

def compare_distributions(low_df, high_df, feature, xlabel, title):
    plt.figure(figsize=(10, 6))
    sns.kdeplot(low_df[feature], fill=True, color='red', label='Tip < 10%')
    sns.kdeplot(high_df[feature], fill=True, color='green', label='Tip > 25%')
    plt.title(title)
    plt.xlabel(xlabel)
    plt.ylabel('Density')
    plt.legend()
    plt.show()

compare_distributions(low_tip_df, high_tip_df, 'trip_distance',
                    'Trip Distance (miles)', 'Comparison of Trip Distance (Low Tip vs. High Tip)')

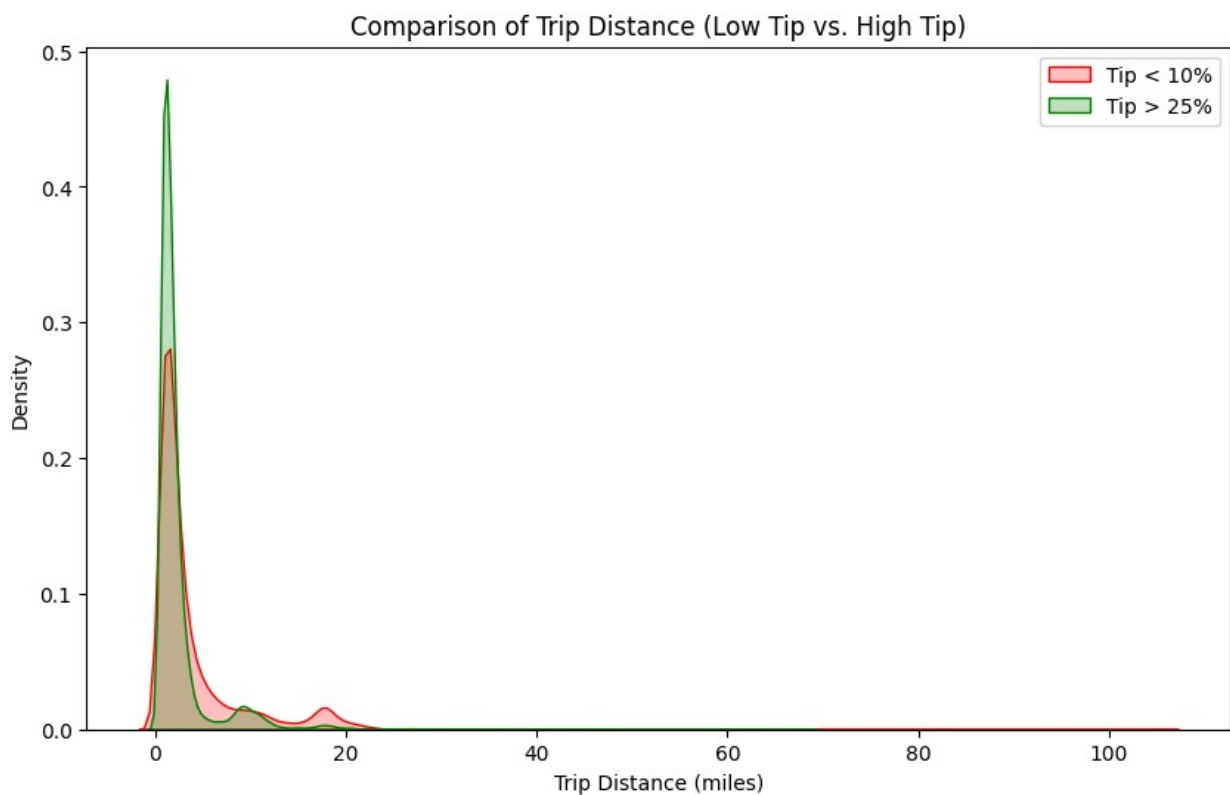
compare_distributions(low_tip_df, high_tip_df, 'fare_amount',
                    'Fare Amount (USD)', 'Comparison of Fare Amount (Low Tip vs. High Tip)')

plt.figure(figsize=(8, 6))
sns.countplot(x='passenger_count', hue=(valid_df['tip_percentage'] > 25), data=valid_df, palette='Set2')
plt.title('Passenger Count Comparison (Low Tip vs. High Tip)')
plt.xlabel('Passenger Count')
plt.ylabel('Number of Trips')
plt.legend(['Tip < 10%', 'Tip > 25%'])
```


```
plt.show()


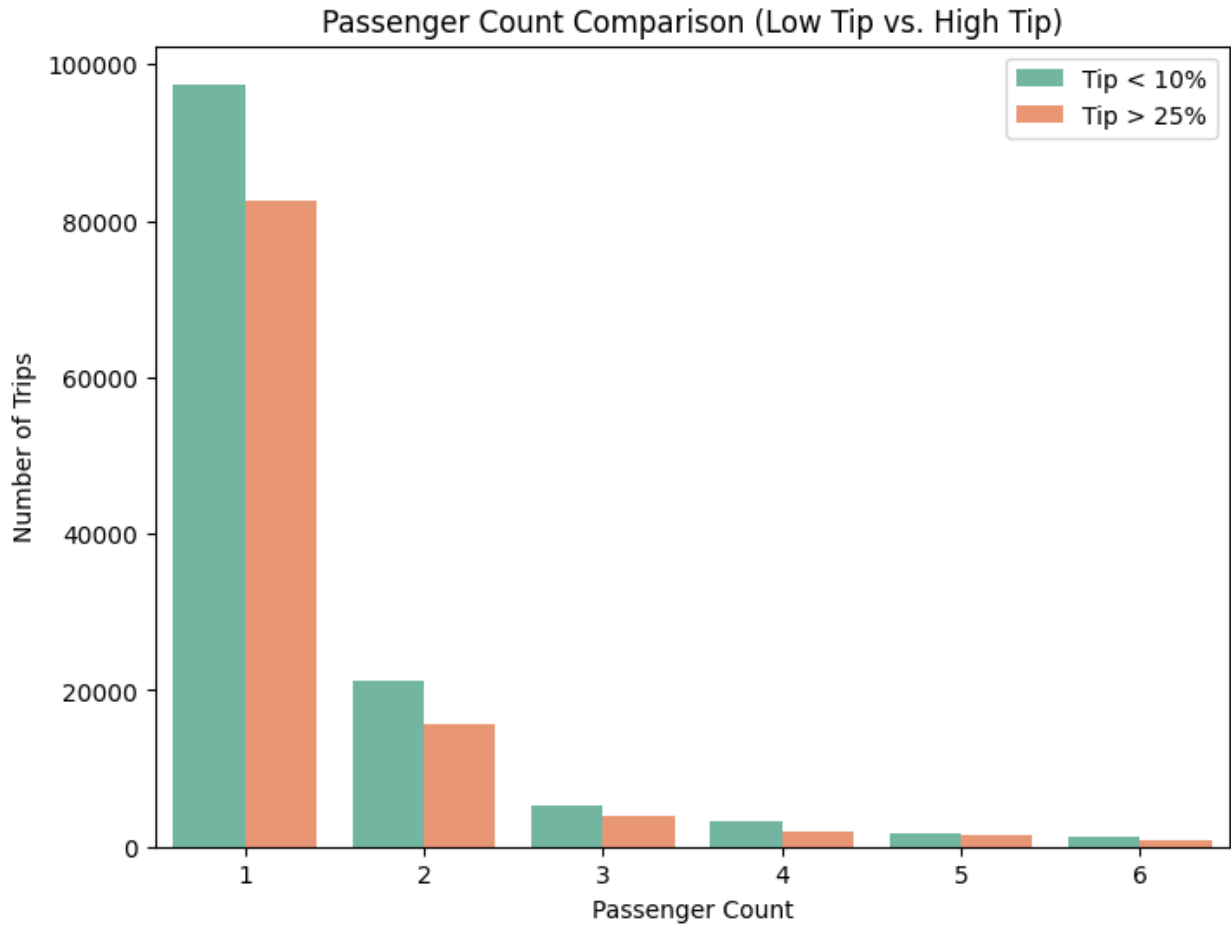
low_tip_df['pickup_hour'] = low_tip_df['tpep_pickup_datetime'].dt.hour
high_tip_df['pickup_hour'] =
high_tip_df['tpep_pickup_datetime'].dt.hour

plt.figure(figsize=(12, 6))
sns.histplot(low_tip_df['pickup_hour'], kde=True, color='red',
label='Tip < 10%', bins=24)
sns.histplot(high_tip_df['pickup_hour'], kde=True, color='green',
label='Tip > 25%', bins=24)
plt.title('Comparison of Pickup Times (Low Tip vs. High Tip)')
plt.xlabel('Hour of Day (0-23)')
plt.ylabel('Trip Frequency')
plt.legend()
plt.show()
```



Comparison of Fare Amount (Low Tip vs. High Tip)





3.2.14 [3 marks] Analyse the variation of passenger count across hours and days of the week.

```
# See how passenger count varies across hours and days
```

```
df['tpep_pickup_datetime'] =  
pd.to_datetime(df['tpep_pickup_datetime'])
```


```
df['pickup_hour'] = df['tpep_pickup_datetime'].dt.hour  
df['pickup_day'] = df['tpep_pickup_datetime'].dt.dayofweek
```

```
day_map = {0: 'Monday', 1: 'Tuesday', 2: 'Wednesday', 3: 'Thursday',  
4: 'Friday', 5: 'Saturday', 6: 'Sunday'}  
df['pickup_day_name'] = df['pickup_day'].map(day_map)
```

```
pivot_table = df.pivot_table(index='pickup_day_name',  
columns='pickup_hour', values='passenger_count', aggfunc='mean')
```

```
pivot_table = pivot_table.loc[list(day_map.values())]
```

```
plt.figure(figsize=(14, 8))  
sns.heatmap(pivot_table, cmap='YlOrRd', annot=True, fmt=".2f",  
linewidths=0.5)  
plt.title('Average Passenger Count by Hour and Day of the Week')  
plt.xlabel('Hour of the Day (0-23)')  
plt.ylabel('Day of the Week')  
plt.show()
```

3.2.15 [2 marks] Analyse the variation of passenger counts across zones

How does passenger count vary across zones

```

zone_counts = df['PULocationID'].value_counts()
top_zones = zone_counts[zone_counts > 50].index # Filter zones with
more than 50 trips
filtered_df = df[df['PULocationID'].isin(top_zones)]

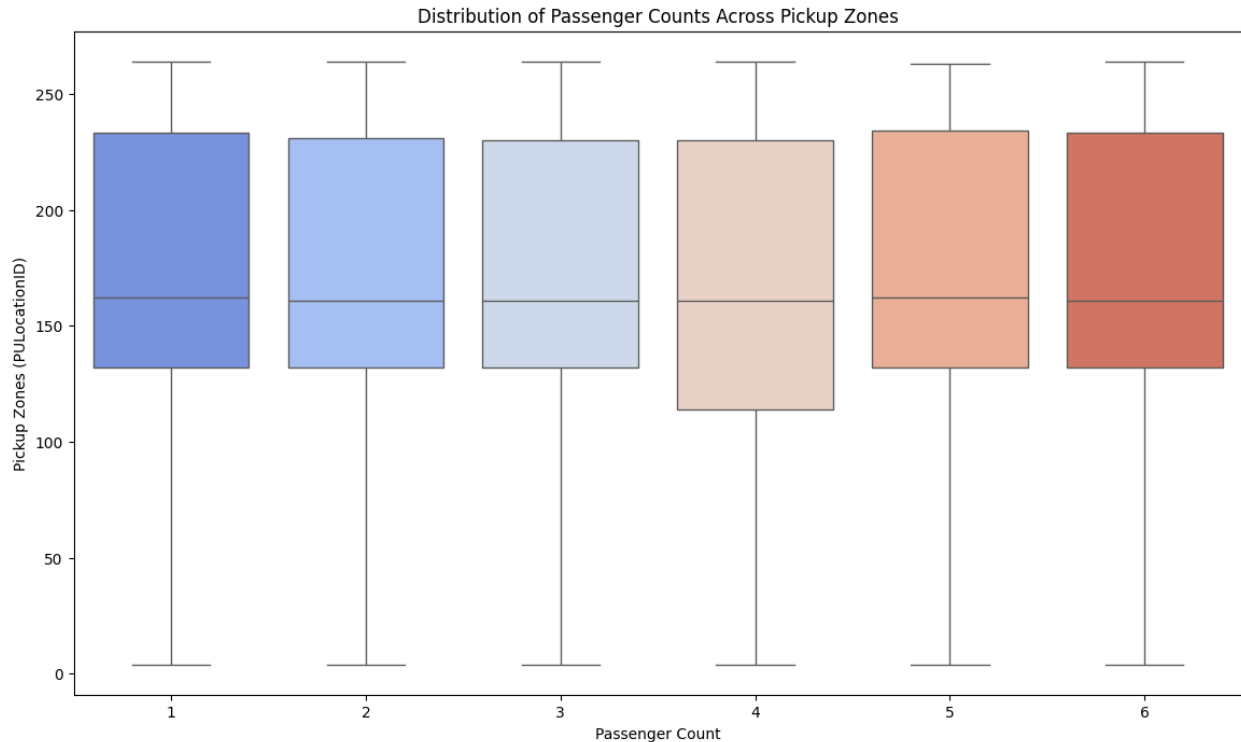
plt.figure(figsize=(14, 8))

sns.boxplot(x='passenger_count', y='PULocationID', data=filtered_df,
palette='coolwarm', showfliers=False)

plt.title('Distribution of Passenger Counts Across Pickup Zones')
plt.xlabel('Passenger Count')
plt.ylabel('Pickup Zones (PULocationID)')

plt.show()

```



```
# For a more detailed analysis, we can use the zones_with_trips
GeoDataFrame
# Create a new column for the average passenger count in each zone.
```

Find out how often surcharges/extra charges are applied to understand their prevalence

3.2.16 [5 marks] Analyse the pickup/dropoff zones or times when extra charges are applied more frequently

```
# How often is each surcharge applied?
```

```
surcharge_columns = ['extra', 'mta_tax', 'tip_amount', 'tolls_amount',
                     'improvement_surcharge', 'congestion_surcharge',
                     'Combined_Airport_Fee']
```

```
surcharge_counts = (df[surcharge_columns] != 0).sum().reset_index()
surcharge_counts.columns = ['Surcharge', 'Frequency']
```

```
print(surcharge_counts)
```

	Surcharge	Frequency
0	extra	146570
1	mta_tax	235047

2	tip_amount	186951
3	tolls_amount	19980
4	improvement_surcharge	236457
5	congestion_surcharge	221655
6	Combined_Airport_Fee	21824

4 Conclusion

[15 marks]

4.1 Final Insights and Recommendations

[15 marks]

Conclude your analyses here. Include all the outcomes you found based on the analysis.

Based on the insights, frame a concluding story explaining suitable parameters such as location, time of the day, day of the week etc. to be kept in mind while devising a strategy to meet customer demand and optimise supply.

4.1.1 [5 marks] Recommendations to optimize routing and dispatching based on demand patterns and operational inefficiencies

```
optimization_strategy = """
Dynamic Dispatching:
    - Prioritize high-demand zones during peak hours (morning rush,
evening, night surge).
    - Use real-time monitoring to adjust vehicle allocation.

Geofencing & Supply Management:
    - Set up geofences around busy areas (airports, business hubs).
    - Move idle vehicles to under-served regions.

Route Optimization:
    - Use algorithms for shortest, most efficient routes.
    - Minimize empty return trips by smart fleet positioning.

Driver Incentives:
    - Offer bonuses for low-demand areas and off-peak hours.
    - Reward long-distance trips to ensure broader coverage.
"""
```

```
print(optimization_strategy)
```

```
Dynamic Dispatching:
    - Prioritize high-demand zones during peak hours (morning rush,
evening, night surge).
    - Use real-time monitoring to adjust vehicle allocation.
```

Geofencing & Supply Management:

- Set up geofences around busy areas (airports, business hubs).
- Move idle vehicles to under-served regions.

Route Optimization:

- Use algorithms for shortest, most efficient routes.
- Minimize empty return trips by smart fleet positioning.

Driver Incentives:

- Offer bonuses for low-demand areas and off-peak hours.
- Reward long-distance trips to ensure broader coverage.

4.1.2 [5 marks]

Suggestions on strategically positioning cabs across different zones to make best use of insights uncovered by analysing trip trends across time, days and months.

```
cab_strategy = """
Peak Hour Hotspots:
    - Morning (7 AM - 10 AM): Position cabs in residential zones for
work commutes.
    - Evening (5 PM - 8 PM): Focus on business districts for return
trips.

Night Demand Zones:
    - 11 PM - 3 AM: Place cabs near entertainment hubs, airports, and
transport stations to capture late-night travelers.

Weekend Strategy:
    - Increase cab availability in malls, tourist spots, and event
venues during weekends (especially in the afternoon and evening).

Low-Demand Redistribution:
    - Reallocate idle cabs from low-traffic areas to zones with
emerging trends (e.g., newly developed regions or seasonal event
areas).

Seasonal Adjustment:
    - Adjust cab supply based on monthly trends—increase near holiday
destinations during vacation seasons and business hubs during work
periods.

Zone-Specific Supply Balancing:
    - Ensure a balanced distribution by analyzing pickup/drop-off
imbalances—position more cabs in zones with high outbound demand.
"""

print(cab_strategy)
```

Peak Hour Hotspots:

- Morning (7 AM - 10 AM): Position cabs in residential zones for work commutes.
- Evening (5 PM - 8 PM): Focus on business districts for return trips.

Night Demand Zones:

- 11 PM - 3 AM: Place cabs near entertainment hubs, airports, and transport stations to capture late-night travelers.

Weekend Strategy:

- Increase cab availability in malls, tourist spots, and event venues during weekends (especially in the afternoon and evening).

Low-Demand Redistribution:

- Reallocate idle cabs from low-traffic areas to zones with emerging trends (e.g., newly developed regions or seasonal event areas).

Seasonal Adjustment:

- Adjust cab supply based on monthly trends—increase near holiday destinations during vacation seasons and business hubs during work periods.

Zone-Specific Supply Balancing:

- Ensure a balanced distribution by analyzing pickup/drop-off imbalances—position more cabs in zones with high outbound demand.

4.1.3 [5 marks] Propose data-driven adjustments to the pricing strategy to maximize revenue while maintaining competitive rates with other vendors.

```
pricing_strategy = ""
```

1. Dynamic Surge Pricing:

- Increase fares during peak hours and in high-demand zones.

2. Distance-Based Tiers:

- Lower fares for short trips, standard for medium trips, and discounts for long trips.

3. Passenger-Based Rates:

- Higher rates for 1-2 passengers, discounts for 3+ passengers.

4. Night & Off-Peak Discounts:

- Offer lower fares during off-peak hours (10 PM - 6 AM) to boost demand.

5. Loyalty Programs:

- Introduce discounts for frequent riders and subscription plans

```
for regular commuters.
```

```
6. Competitor Matching:
```

```
- Adjust fares based on competitor pricing to stay competitive.  
"""
```

```
print(pricing_strategy)
```

```
1. Dynamic Surge Pricing:
```

```
- Increase fares during peak hours and in high-demand zones.
```

```
2. Distance-Based Tiers:
```

```
- Lower fares for short trips, standard for medium trips, and  
discounts for long trips.
```

```
3. Passenger-Based Rates:
```

```
- Higher rates for 1-2 passengers, discounts for 3+ passengers.
```

```
4. Night & Off-Peak Discounts:
```

```
- Offer lower fares during off-peak hours (10 PM - 6 AM) to boost  
demand.
```

```
5. Loyalty Programs:
```

```
- Introduce discounts for frequent riders and subscription plans  
for regular commuters.
```

```
6. Competitor Matching:
```

```
- Adjust fares based on competitor pricing to stay competitive.
```